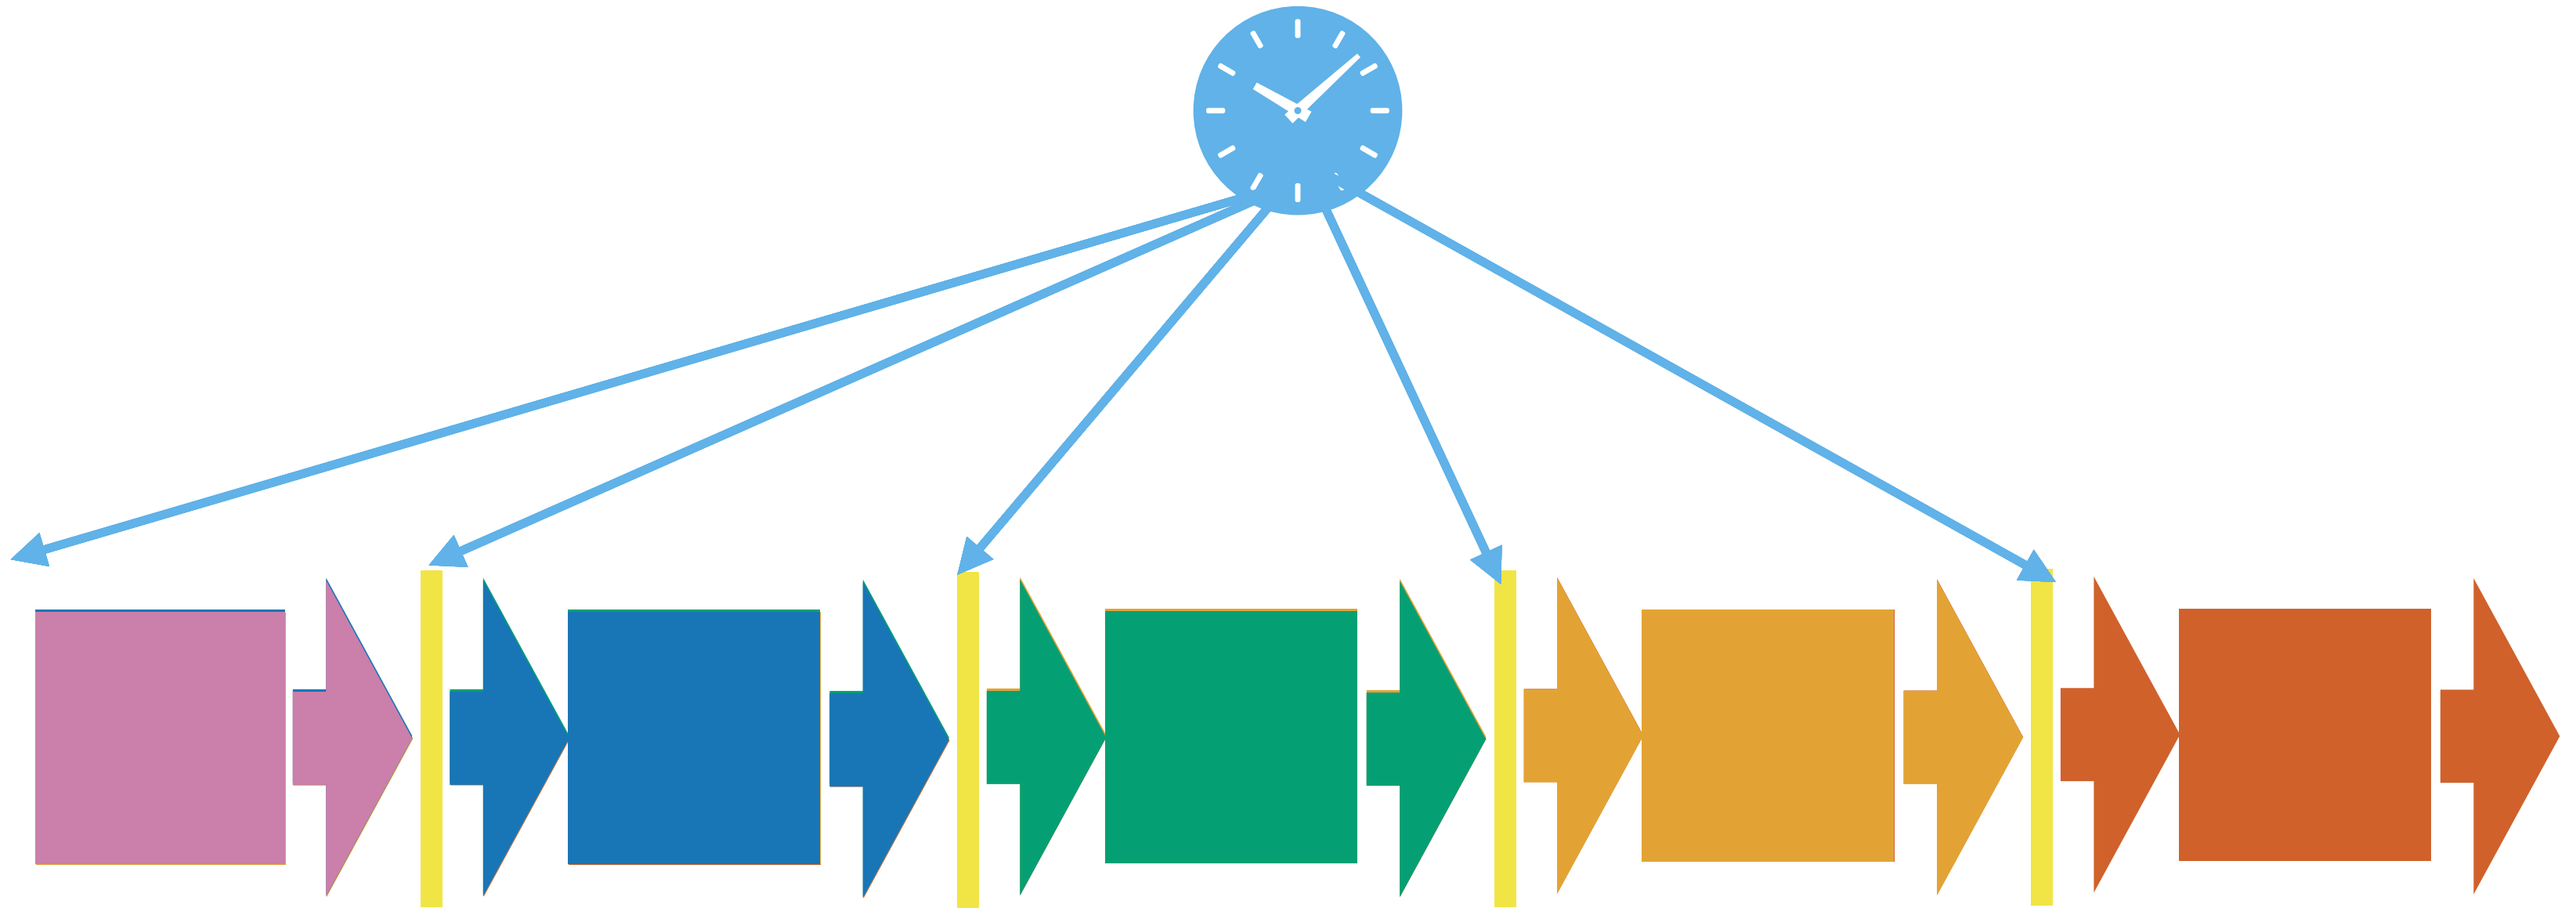


Data Hazards & Dynamic Instruction Scheduling: CPUtopia

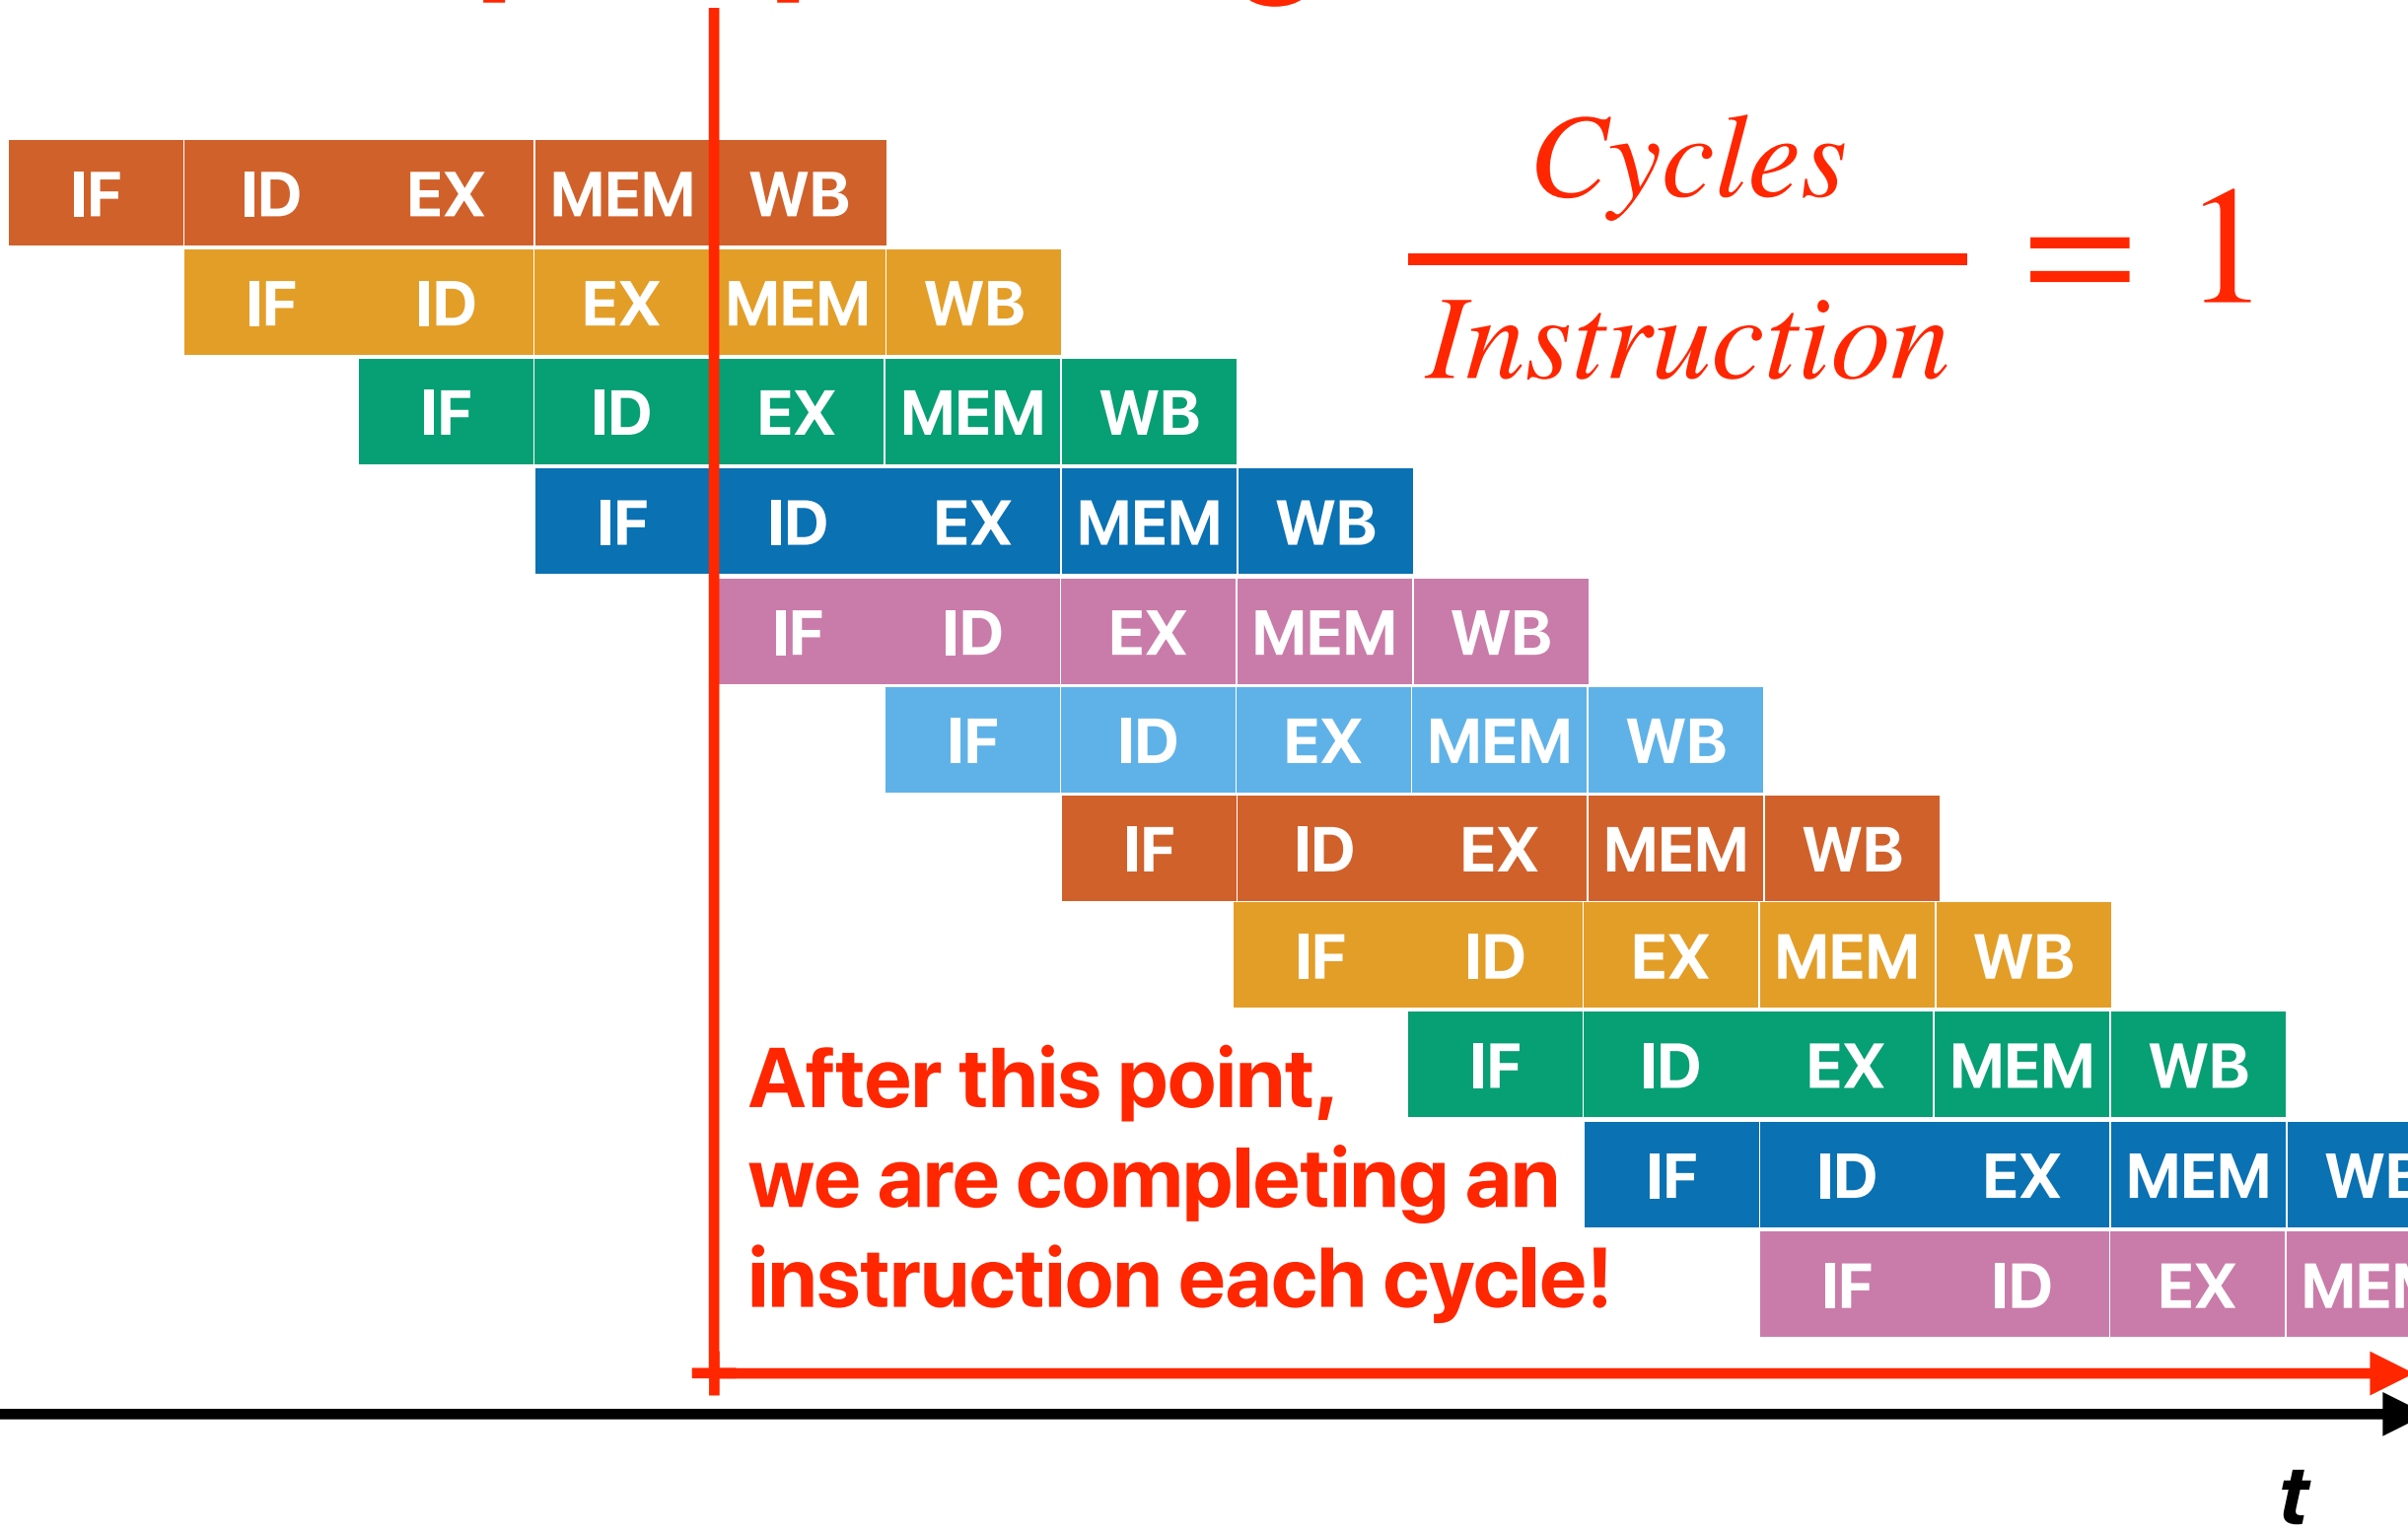
Hung-Wei Tseng

Recap: Pipelining



Recap: Pipelining

```
add x1, x2, x3
ld x4, 0(x5)
sub x6, x7, x8
sub x9, x10, x11
sd x1, 0(x12)
xor x13, x14, x15
and x16, x17, x18
add x19, x20, x21
sub x22, x23, x24
ld x25, 4(x26)
sd x27, 0(x28)
```



Recap: Three pipeline hazards

- Structural hazards — resource conflicts cannot support simultaneous execution of instructions in the pipeline
- Control hazards — the PC can be changed by an instruction in the pipeline
- Data hazards — an instruction depending on a the result that's not yet generated or propagated when the instruction needs that

Recap: addressing hazards

- Structural hazards
 - Stall
 - Modify hardware design
- Control hazards
 - Stall
 - Static prediction
 - Dynamic prediction

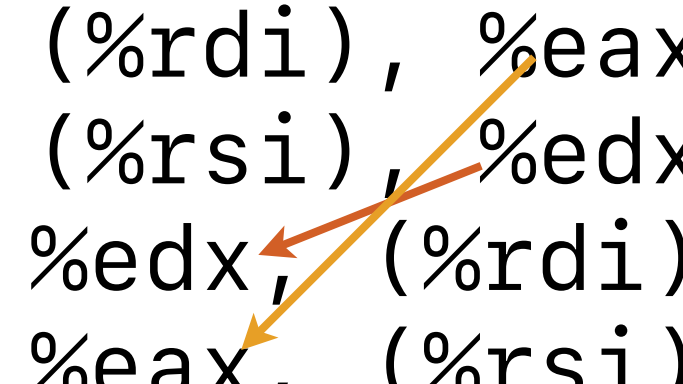
Recap: Data hazards

- An instruction currently in the pipeline cannot receive the “logically” correct value for execution
- Data dependencies
 - The output of an instruction is the input of a later instruction
 - **May (but not necessary)** result in data hazard if the later instruction that consumes the result is still in the pipeline

How many dependencies do we have?

- How many pairs of data dependences are there in the following x86 instructions?

```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
```



```
int temp = *a;
*a = *b;
*b = temp;
```

A. 1

B. 2

C. 3

D. 4

E. 5

How many dependencies do we have?

- How many pairs of data dependencies are there in the following x86 instructions?

```
movl    (%rdi), %eax
xorl    (%rsi), %eax
movl    %eax, (%rdi)
xorl    (%rsi), %eax
movl    %eax, (%rsi)
xorl    %eax, (%rdi)
```

```
*a ^= *b;
*b ^= *a;
*a ^= *b;
```

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5

**What will you do if you're waiting for
someone else's response or
outcome?**

Ideas?

Outline

- Data hazards
 - Data forwarding
- SuperScalar
- Out-of-order, Dynamic instruction scheduling

Solution 1: Let's try "stall" again

- Whenever the input is not ready when the consumer is decoding, just stall — the consumer stays at ID.

① `movl (%rdi), %eax`
② `movl (%rsi), %edx`
③ `movl %edx, (%rdi)`
④ `movl %eax, (%rsi)`

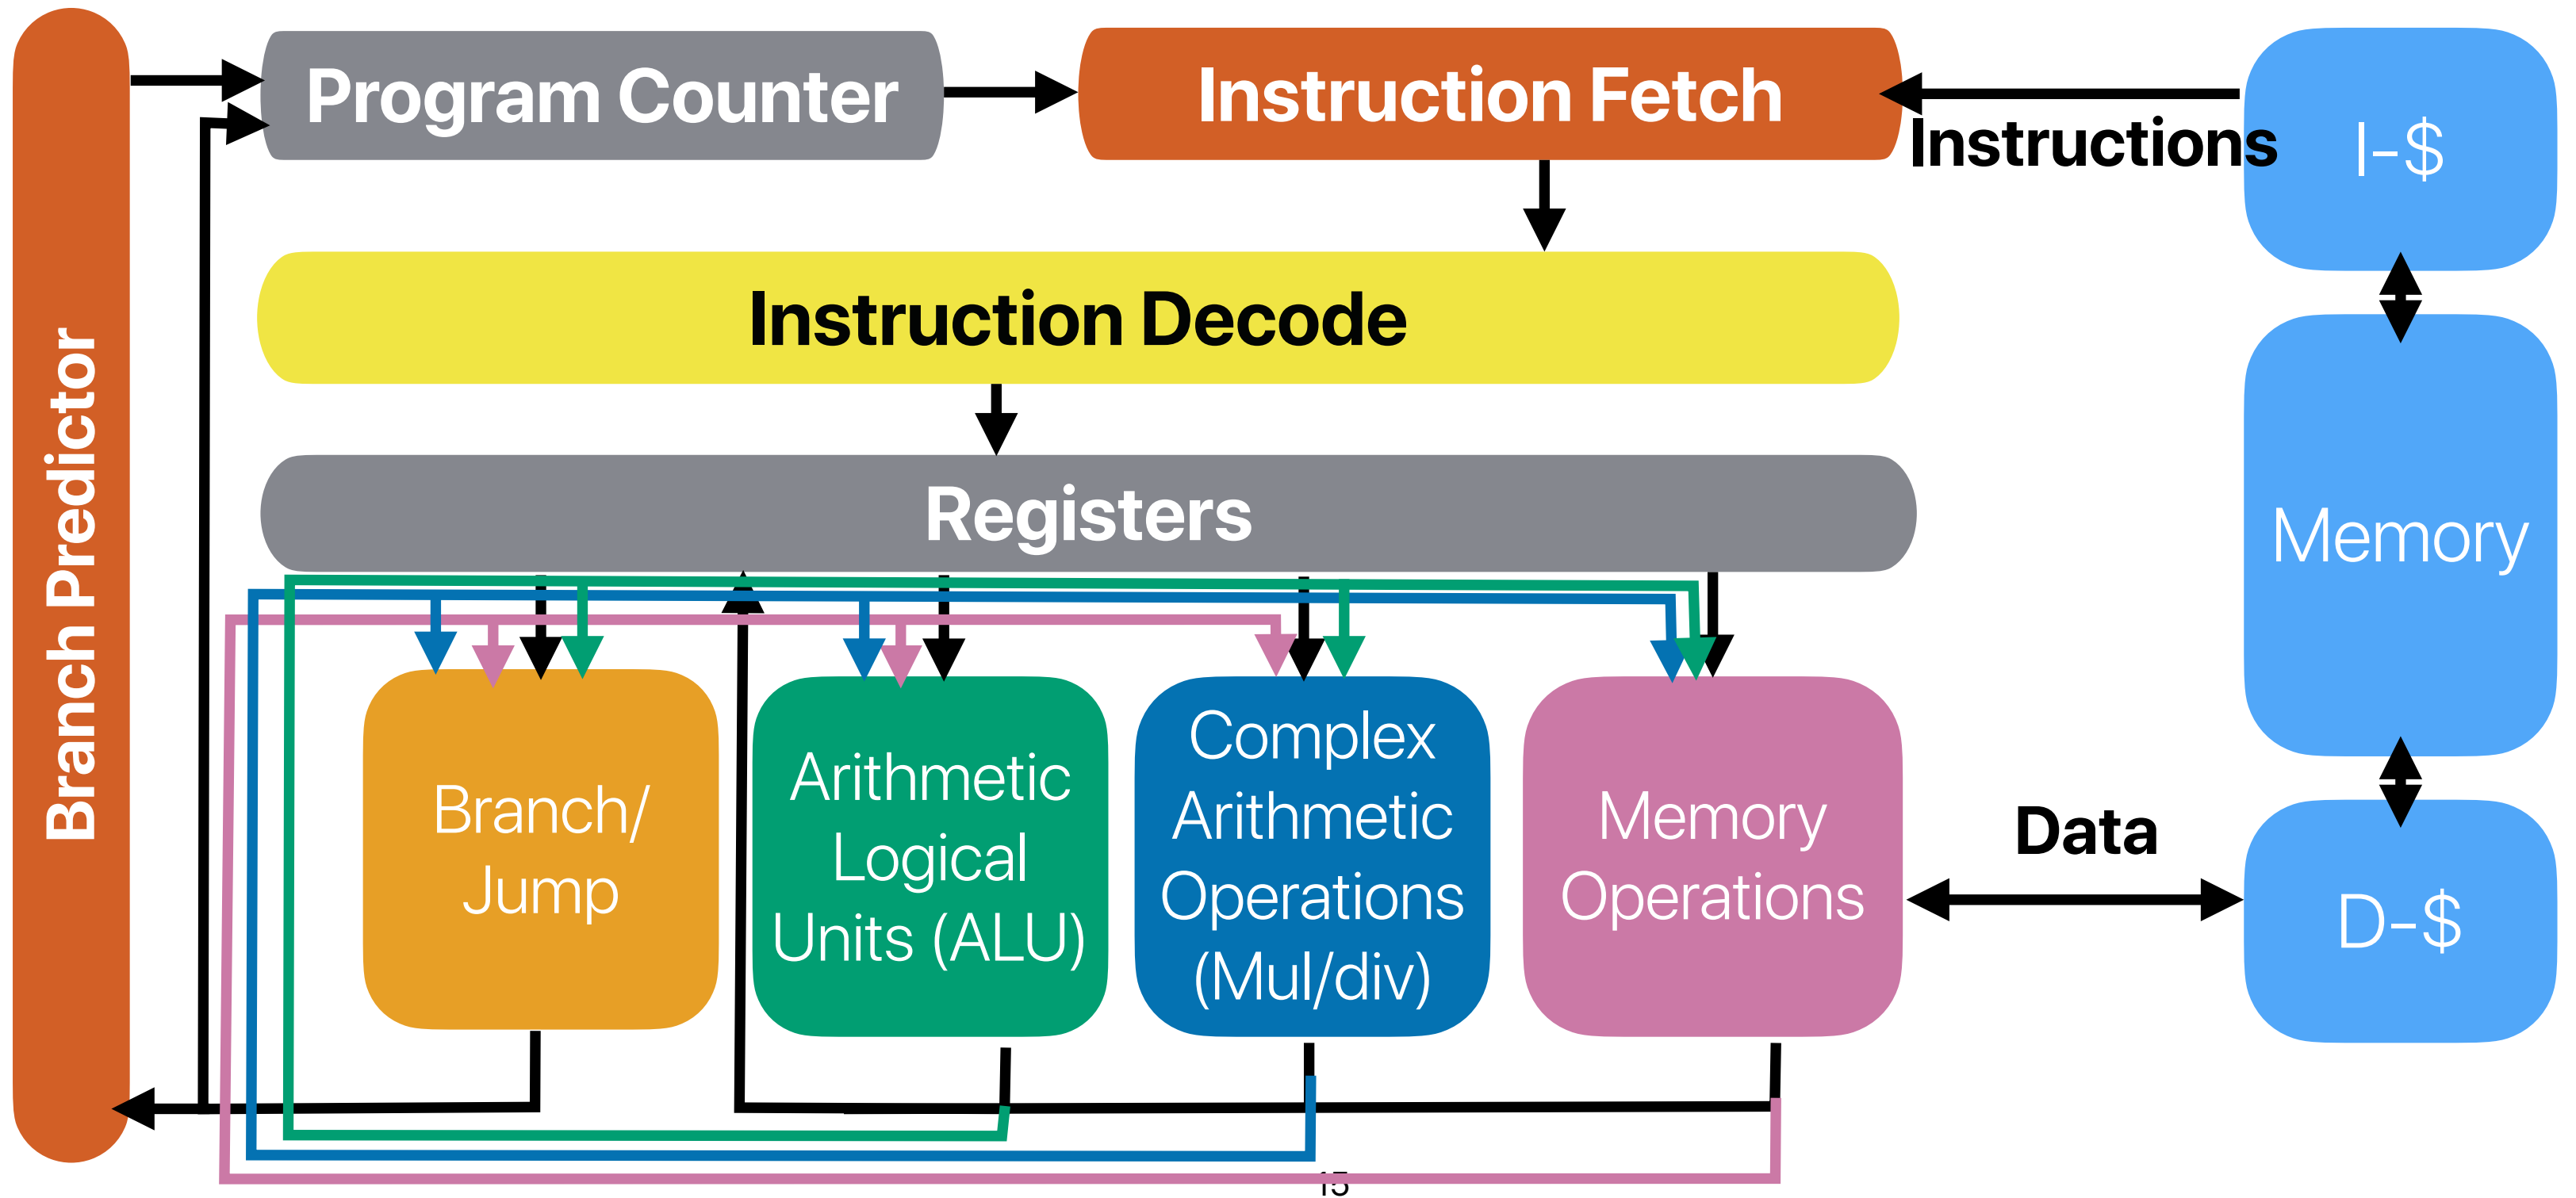
we have the value for %edx already!
Why another cycle?

	IF	ID	ALU/BR/M1	M2	M3	M4	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(4)	(3)	(2)	(1)			
5	(4)	(3)		(2)	(1)		
6	(4)	(3)			(2)	(1)	
7	(4)	(3)				(2)	(1)
8	(4)	(3)					(2)
9		(4)	(3)				
10			(4)	(3)			
11				(4)	(3)		
12					(4)	(3)	
13						(4)	(3)
14							(4)

Solution 2: Data forwarding

- Add logics/wires to forward the desired values to the demanding instructions

Data "forwarding"




How many of them are still problematic?

- How many pairs of data dependences in the following x86 instructions are still problematic with data forwarding if a memory operation takes 4 cycles?

① `movl (%rdi), %eax`
② `movl (%rsi), %edx`
③ `movl %edx, (%rdi)`
④ `movl %eax, (%rsi)`

- A. 0
B. 1
C. 2
D. 3
E. 4



How many of data hazards?

- How many pairs of data dependences in the following x86 instructions are still problematic with data forwarding if a memory operation takes 4 cycles?

① movl (%rdi), %eax

② movl (%rsi), %edx

③ movl %edx, (%rdi)

④ movl %eax, (%rsi)

A. 0

B. 1

C. 2

D. 3

E. 4

```
int temp = *a;  
*a = *b;  
*b = temp;
```

	IF	ID	ALU/BR/M1	M2	M3	M4	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(4)	(3)	(2)	(1)			
5	(4)	(3)		(2)	(1)		
6	(4)	(3)			(2)	(1)	
7		(4)				(2)	(1)
8			(3)				(2)
9			(4)	(3)			
10				(4)	(3)		
11					(4)	(3)	
12						(4)	(3)
13							(4)
14							

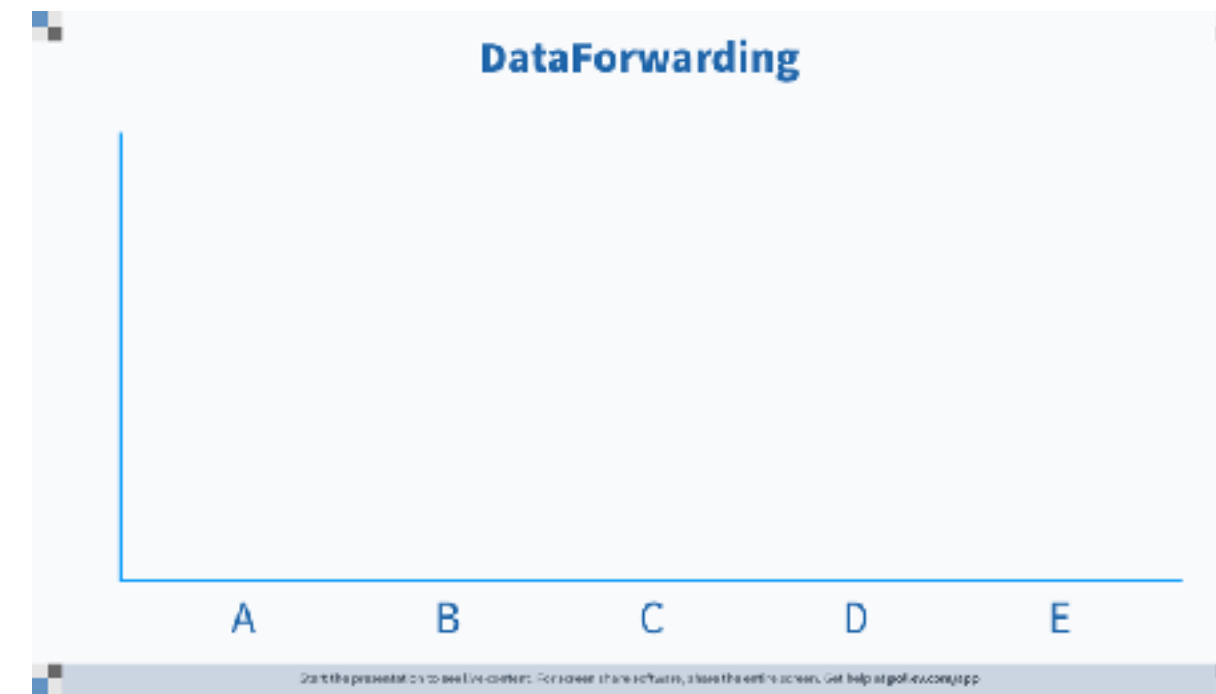
How many of them are still problematic?

- How many pairs of data dependences in the following x86 instructions are still problematic with data forwarding if a memory operation takes 4 cycles and **xorl also takes 4 cycles**?

① `movl (%rdi), %eax`
② `xorl (%rsi), %eax`
③ `movl %eax, (%rdi)`
④ `xorl (%rsi), %eax`
⑤ `movl %eax, (%rsi)`
⑥ `xorl %eax, (%rdi)`

`*a ^= *b;`
`*b ^= *a;`
`*a ^= *b;`

- A. 0
B. 1
C. 2
D. 3
E. 4



How many of data hazards w/ Data Forwarding?

- How many pairs of data dependences in the following x86 instructions are still problematic with data forwarding if a memory operation takes 4 cycles and **xorl also takes 4 cycles**?

① movl (%rdi), %eax

② xorl (%rsi), %eax

③ movl %eax, (%rdi)

④ xorl (%rsi), %eax

⑤ movl %eax, (%rsi)

⑥ xorl %eax, (%rdi)

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4

	IF	ID	ALU/BR/M1	M2	M3	M4/XORL	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(3)	(2)		(1)			
5	(3)	(2)			(1)		
6	(3)	(2)				(1)	
7	(4)	(3)	(2)				(1)
8	(4)	(3)		(2)			
9	(4)	(3)			(2)		
10	(4)	(3)				(2)	
11	(5)	(4)	(3)				(2)
12	(6)	(5)	(4)	(3)			
13	(6)	(5)		(4)	(3)		
14	(6)	(5)			(4)	(3)	
15	(6)	(5)				(4)	(3)
16		(6)	(5)				(4)
17			(6)	(5)			
18				(6)	(5)		
19					(6)	(5)	
20						(6)	(5)
21							(6)
22							

Single pipeline

```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```

```
.L3:  
①    movl    (%rdi), %ecx  
②    addl    %ecx, %eax  
③    addq    $4, %rdi  
④    cmpq    %rdx, %rdi  
⑤    jne     .L3  
⑥    ret
```

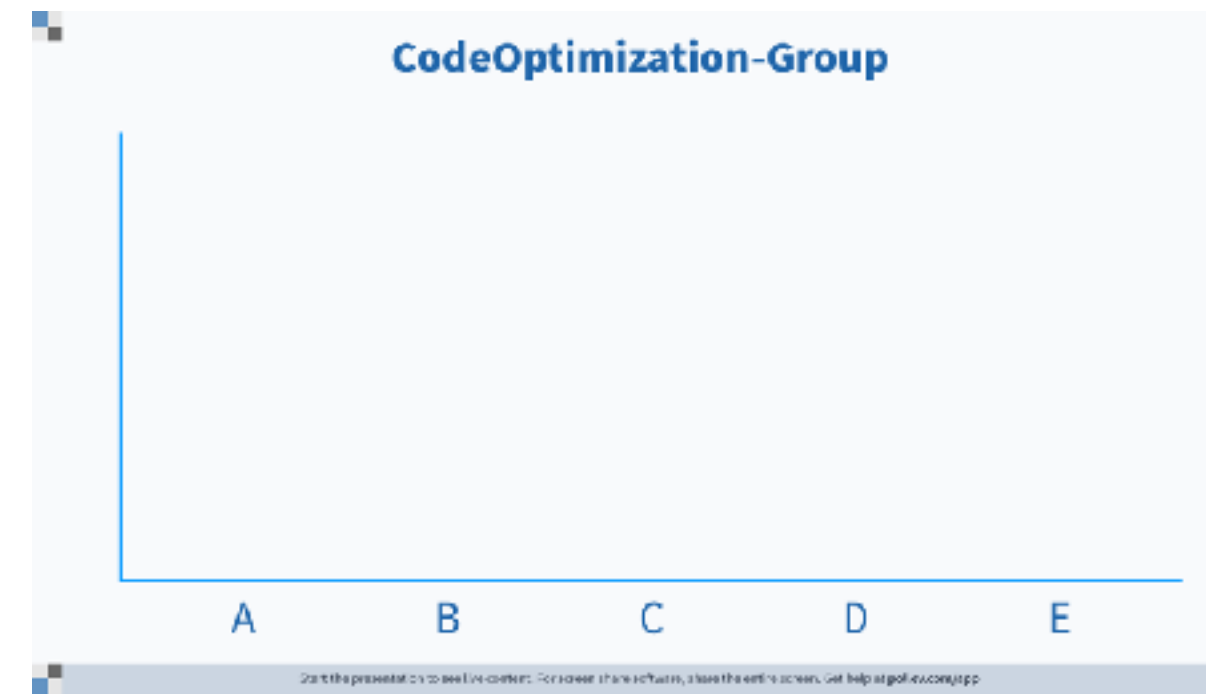
	IF	ID	ALU/BR/M1	M2	M3	M4/XORL	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(3)	(2)		(1)			
5	(3)	(2)			(1)		
6	(3)	(2)				(1)	
7	(4)	(3)	(2)				(1)
8	(5)	(4)	(3)	(2)			
9		(5)	(4)	(3)	(2)		
10			(5)	(4)	(3)	(2)	
11				(5)	(4)	(3)	(2)
12					(5)	(4)	(3)
13						(5)	(4)

The effect of code optimization

- By reordering which pair of the following instruction stream can we not affect the correctness of the code?

- ① `movl (%rdi), %ecx`
- ② `addl %ecx, %eax`
- ③ `addq $4, %rdi`
- ④ `cmpq %rdx, %rdi`
- ⑤ `jne .L3`
- ⑥ `ret`

- A. (1) & (2)
- B. (2) & (3)
- C. (3) & (4)
- D. (4) & (5)
- E. None of the pairs can be reordered



The effect of code optimization

- By reordering which pair of the following instruction stream can we eliminate all stalls without affecting the correctness of the code?

① `movl (%rdi), %ecx`
② `addl %ecx, %eax`
③ `addq $4, %rdi`
④ `cmpq %rdx, %rdi`
⑤ `jne .L3`
⑥ `ret`

A. (1) & (2)

B. (2) & (3)

C. (3) & (4)

D. (4) & (5)

E. None of the pairs can be reordered

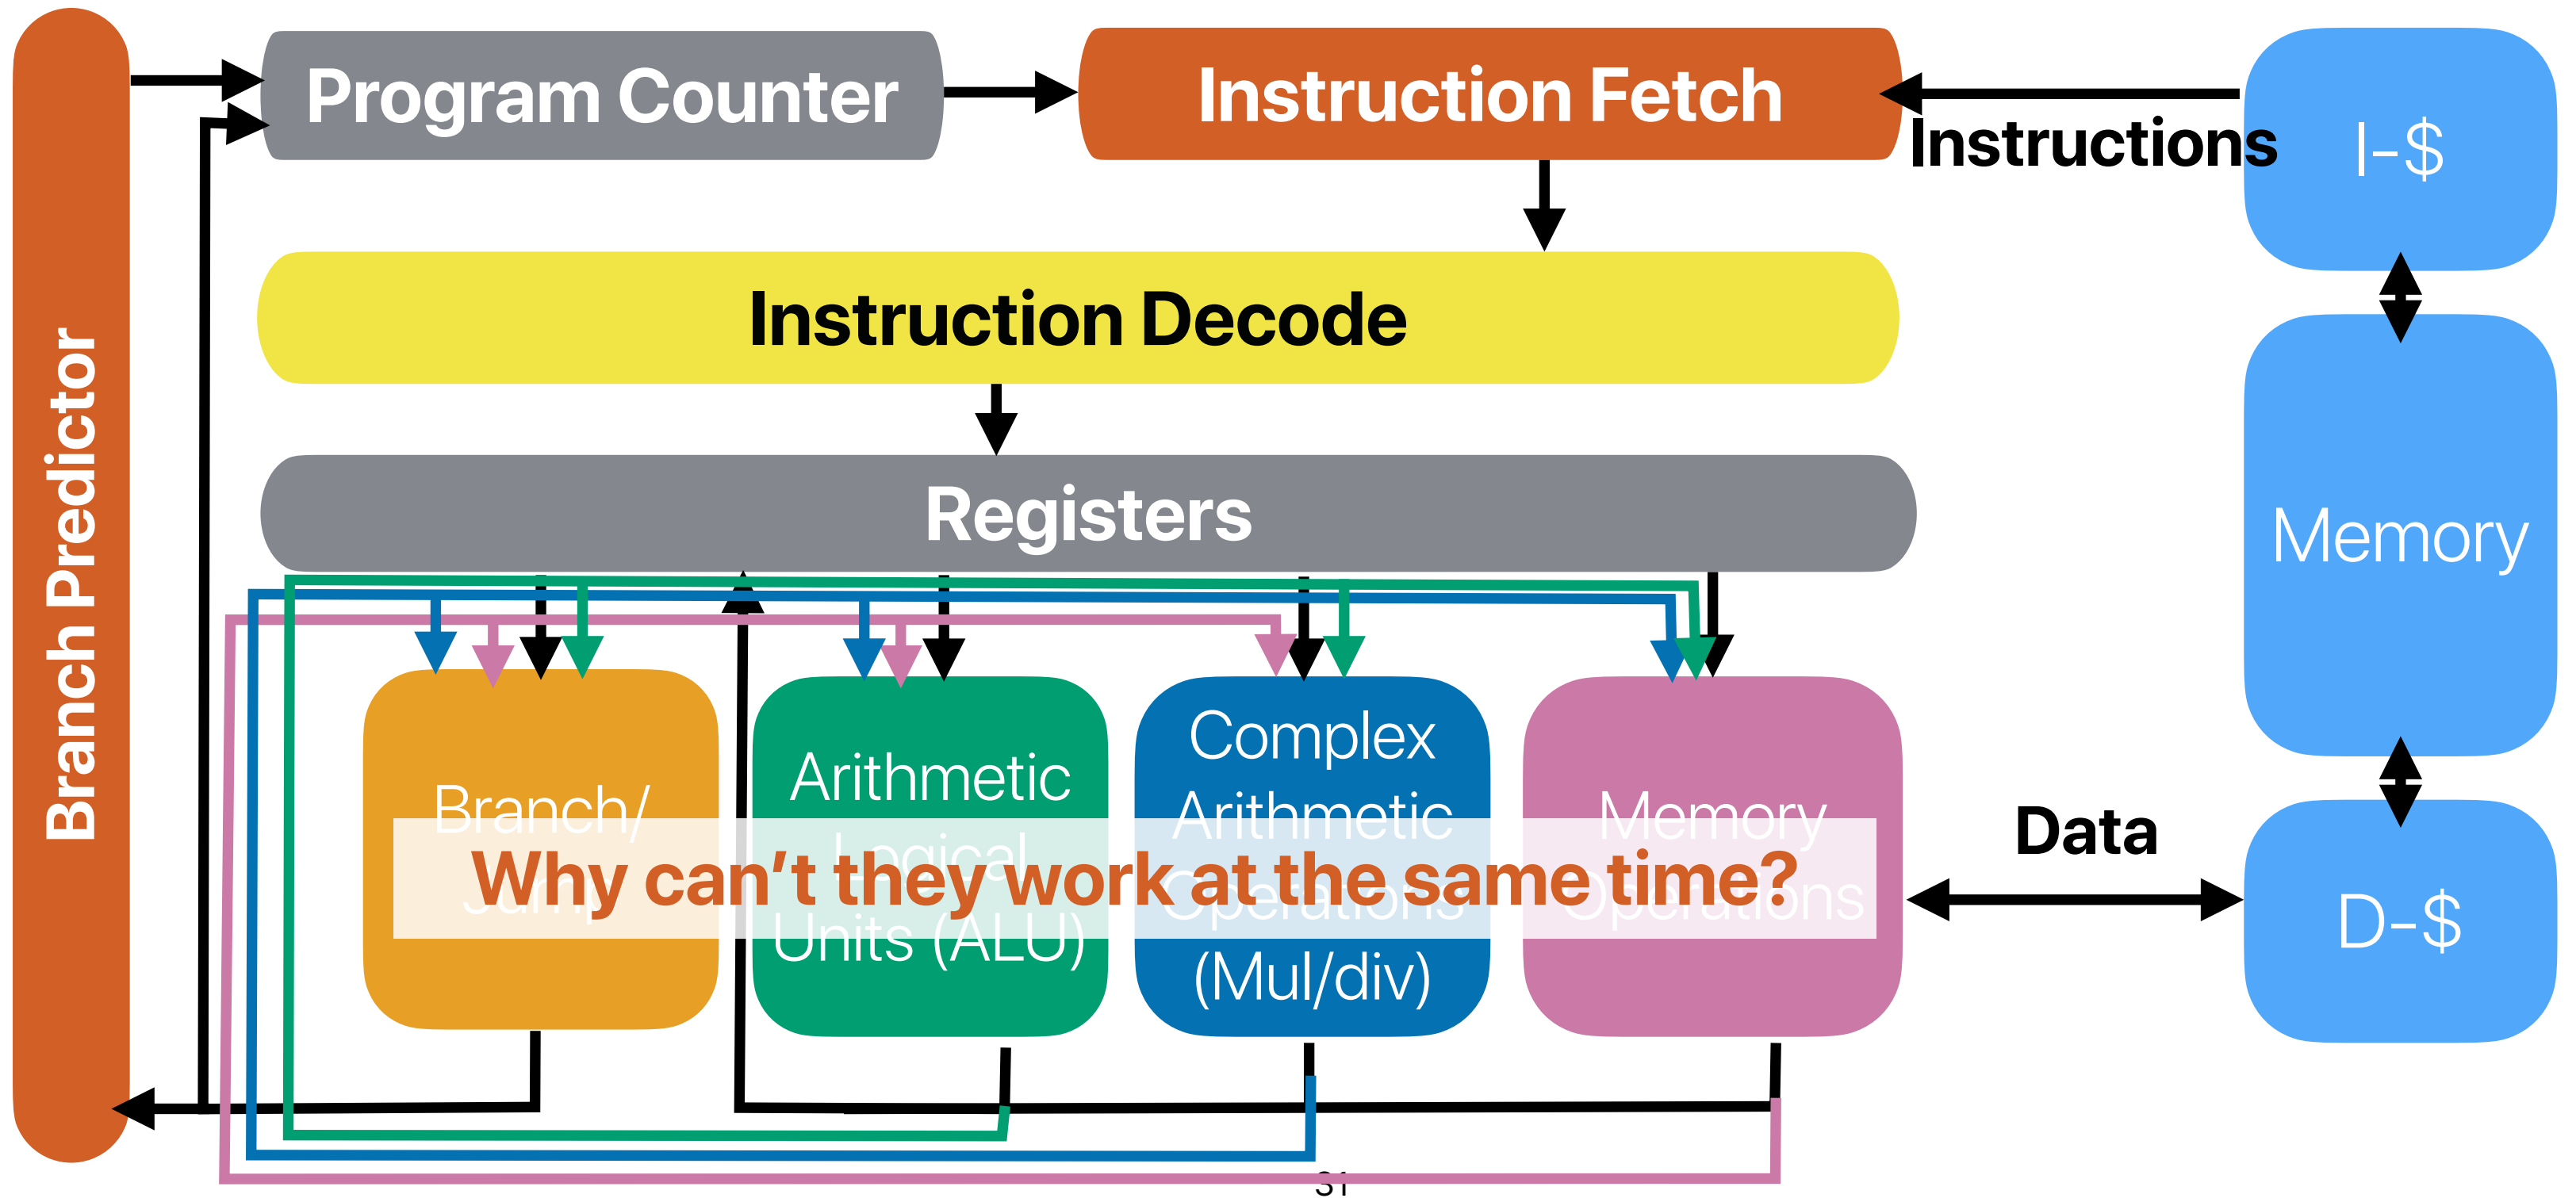
Compiler optimization

```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```

```
.L3:  
①    movl    (%rdi), %ecx  
②    addq    $4, %rdi  
③    addl    %ecx, %eax  
④    cmpq    %rdx, %rdi  
⑤    jne     .L3  
⑥    ret
```

	IF	ID	ALU/BR/M1	M2	M3	M4/XORL	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(3)	(2)	(2)	(1)			
5	(3)	(2)		(2)	(1)		
6	(4)	(2)			(2)	(1)	
7	(5)	(4)	(3)			(2)	(1)
8		(5)	(4)	(3)			(2)
9			(5)	(4)	(3)		
10				(5)	(4)	(3)	
11					(5)	(4)	(3)
12						(5)	(4)
13							(5)

Data "forwarding"



Compiler optimization

```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```

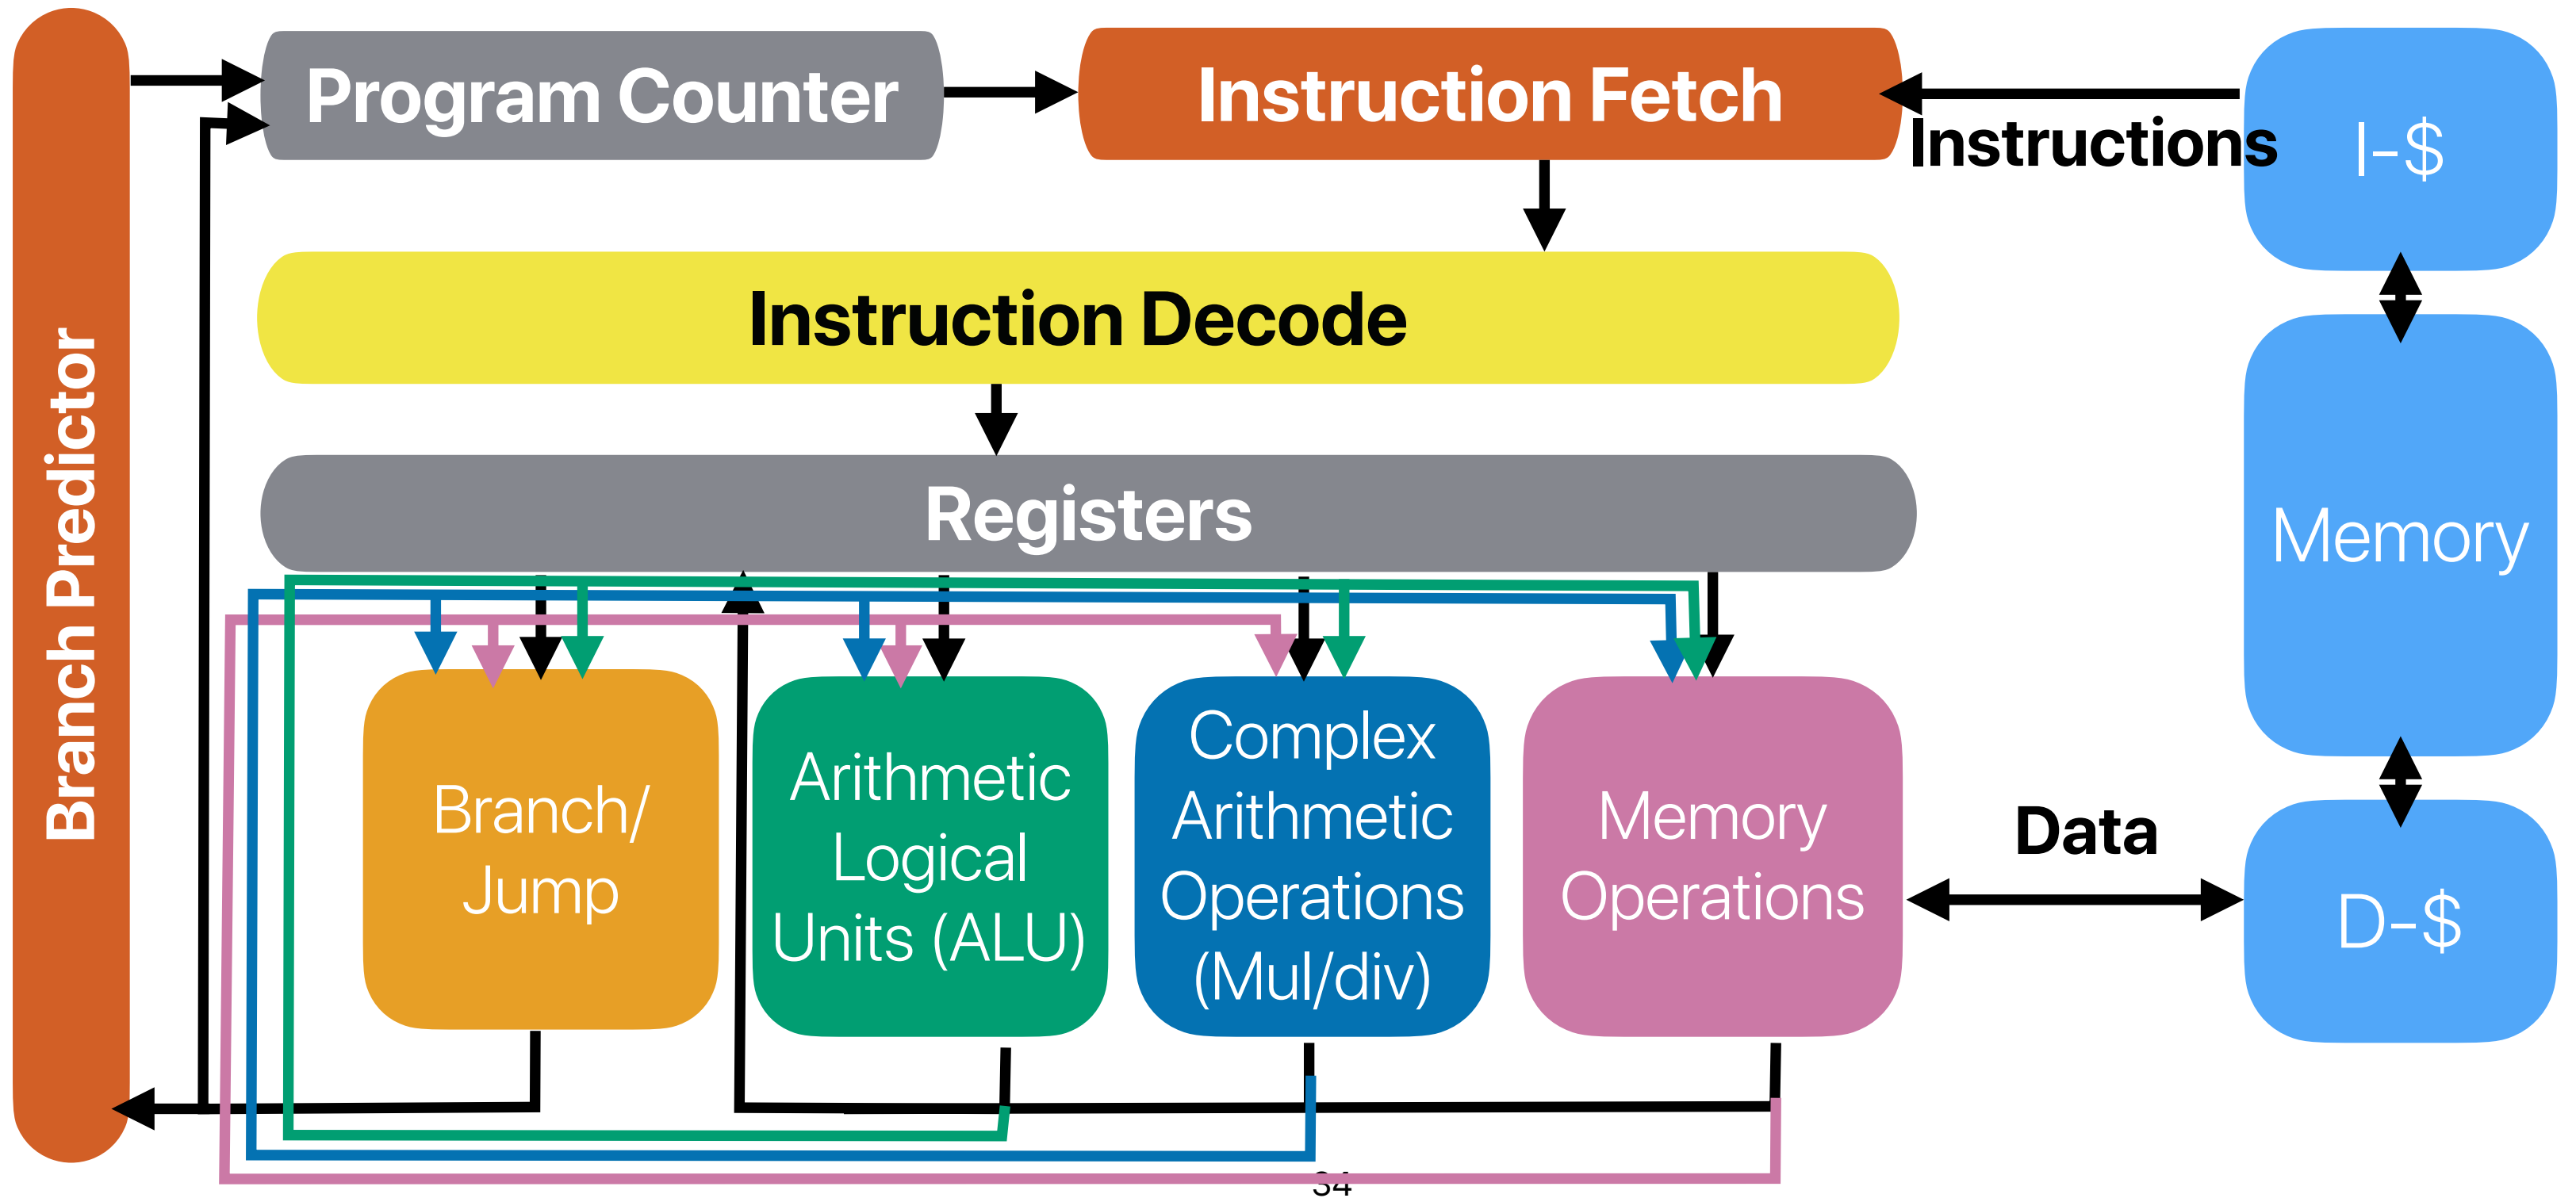
```
.L3:  
①    movl    (%rdi), %ecx  
②    addq    $4, %rdi  
③    addl    %ecx, %eax  
④    cmpq    %rdx, %rdi  
⑤    jne     .L3  
⑥    ret
```

	IF	ID	ALU/BR/M1	M2	M3	M4/XORL	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(3)	(2)	(2)	(1)			
5	(3)	(2)		(2)	(1)		
6	(4)	(2)			(2)	(1)	
7	(5)	(4)	(3)			(2)	(1)
8		(5)	(4)	(3)			(2)
9			(5)	(4)	(3)		
10				(5)	(4)	(3)	
11					(5)	(4)	(3)
12						(5)	(4)
13							(5)

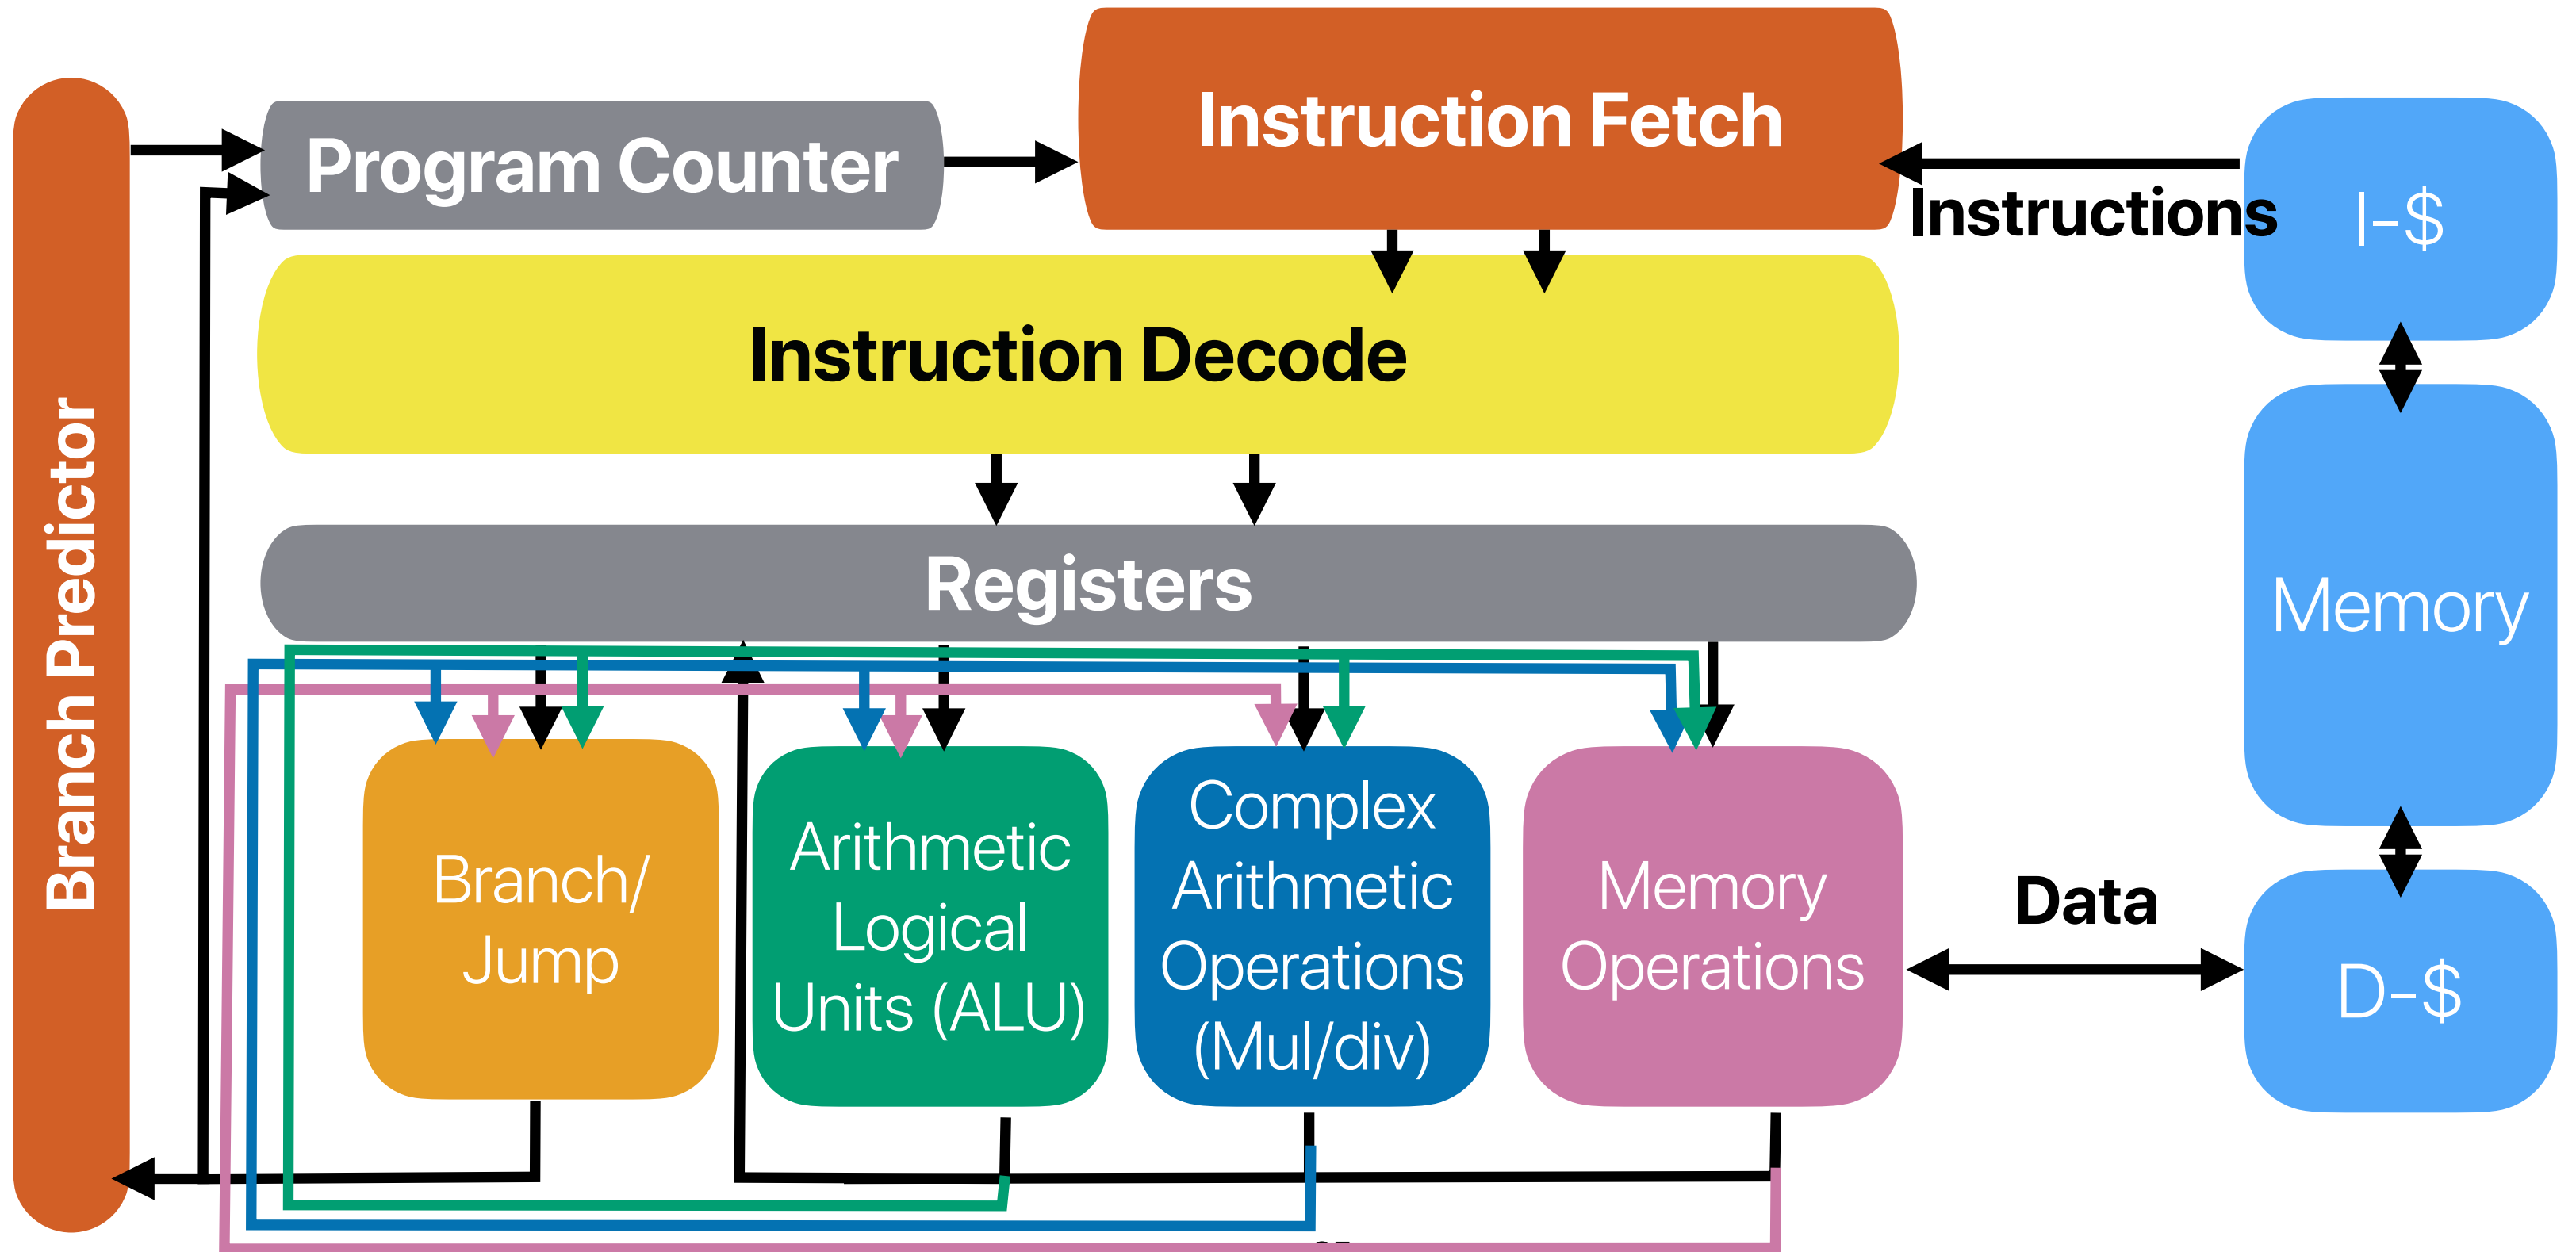
addq is not depending on movl and
ALU is free! can we execute them
together?

If $CPI == 1$ the limitation?

Data "forwarding"



Super Scalar



Super Scalar

Superscalar

- Since we have many functional units now, we should fetch/decode more instructions each cycle so that we can have more instructions to issue!
- Super-scalar: fetch/decode/issue more than one instruction each cycle
 - **Fetch width:** how many instructions can the processor fetch/decode each cycle
 - **Issue width:** how many instructions can the processor issue each cycle
- The theoretical CPI should now be

1

$\min(\text{issue width}, \text{fetch width}, \text{decode width})$

Superscalar: fetch/issue width == 2, theoretical CPI = 0.5

```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```

```
.L3:  
①    movl    (%rdi), %ecx  
②    addq    $4, %rdi  
③    addl    %ecx, %eax  
④    cmpq    %rdx, %rdi  
⑤    jne     .L3  
⑥    ret
```

	IF	ID	M1/ALU/BR	M2	M3	M4	WB
1	(1) (2)						
2	(3)(4)	(1) (2)					
3	(5)	(3)(4)	(1)(2)				
4	(5)	(3)(4)		(1)(2)			
5	(5)	(3)(4)			(1)(2)		
6	(5)	(3)(4)				(1)(2)	
7		(5)	(3)(4)				(1)(2)
8			(5)	(3)(4)			
9				(5)	(3)(4)		
10					(5)	(3)(4)	
11						(5)	(3)(4)
12							(5)

If we loop many times (assume perfect predictor)

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
⑪ movl    (%rdi), %ecx
⑫ addq    $4, %rdi
⑬ addl    %ecx, %eax
⑭ cmpq    %rdx, %rdi
⑮ jne     .L3
```

	IF	ID	M1/ALU/BR	M2	M3	M4	WB
1	(1) (2)						
2	(3) (4)	(1) (2)					
3	(5) (6)	(3) (4)	(1) (2)				
4	(5) (6)	(3) (4)		(1) (2)			
5	(5) (6)	(3) (4)			(1) (2)		
6	(5) (6)	(3) (4)				(1) (2)	
7	(7) (8)	(5) (6)	(3) (4)				(1) (2)
8	(9) (10)	(7) (8)	(5) (6)	(3) (4)			
9	(9) (10)	(8)	(7)	(5) (6)	(3) (4)		
10	(9) (10)	(8)		(7)	(5) (6)	(3) (4)	
11	(9) (10)	(8)			(7)	(5) (6)	(3) (4)
12	(11) (12)	(9) (10)	(8)			(7)	(5) (6)
	(11) (12)	(10)	(9)	(8)			(7)
		(11) (12)	(10)	(9)	(8)		
			(11) (12)	(10)	(9)	(8)	
				(11) (12)	(10)	(9)	(8)
					(11)	(10)	(9)

Everything we need
for (4) is ready here
Why can't we
execute it?

Why can't I start
loading (6) & (11)?

Limitations of Compiler Optimizations

- If the hardware (e.g., pipeline changes), the same compiler optimization may not be that helpful
- The compiler can only optimize on static instructions, but cannot optimize dynamic instruction
 - Compiler cannot predict branches
 - Compiler does not know if cache has the data/instructions

What do you need to execution an instruction?

- Whenever the instruction is decoded — put decoded instruction somewhere
- Whenever the inputs are ready — **all data dependencies are resolved**
- Whenever the target functional unit is available

Dynamic instruction scheduling/ Out-of-order (OoO) execution

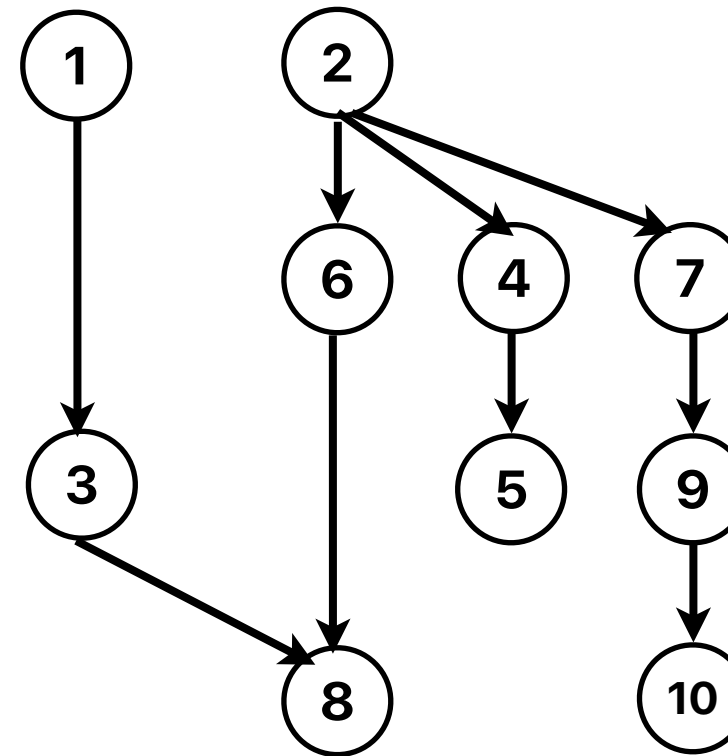
What do you need to execution an instruction?

- Whenever the instruction is decoded — put decoded instruction somewhere
- Whenever the inputs are ready — **all data dependencies are resolved**
- Whenever the target functional unit is available

Scheduling instructions: based on data dependencies

- Draw the data dependency graph, put an arrow if an instruction depends on the other.

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```



- In theory**, instructions without dependencies can be executed in parallel or out-of-order
- Instructions with dependencies can never be reordered

If we can predict the future ...

- Consider the following dynamic instructions:

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```

Which of the following pair can we reorder without affecting the correctness if the **branch prediction is perfect**?

- A. (1) and (2)
- B. (3) and (4)
- C. (3) and (6)
- D. (4) and (7)
- E. (6) and (7)

If we can predict the future ...

- Consider the following dynamic instructions:

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```

Can we use "branch prediction" to predict the future and reorder instructions across the branch?

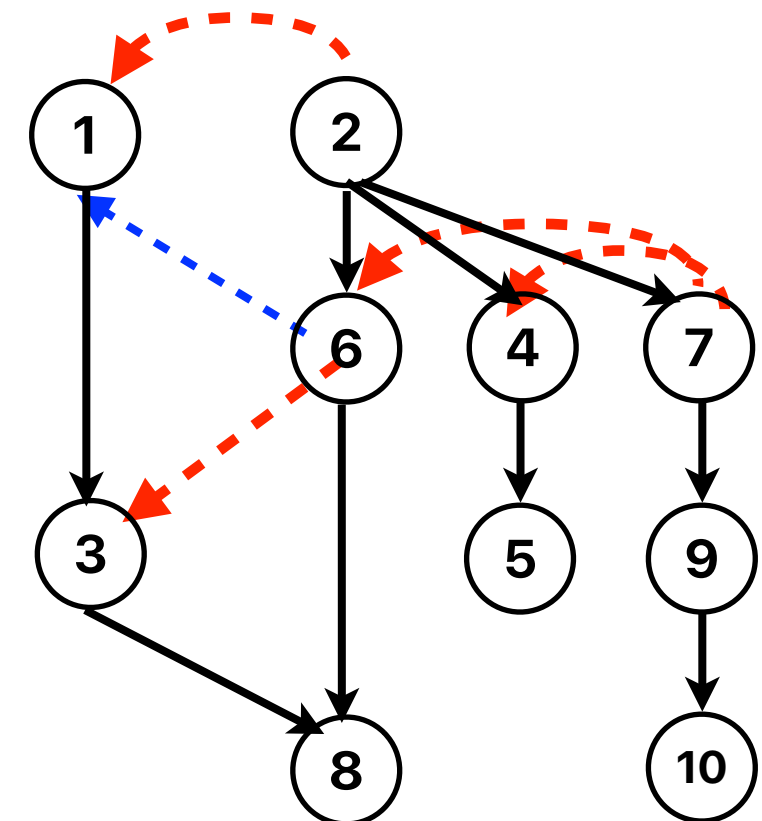
Which of the following pair can we reorder without affecting the correctness if the **branch prediction is perfect**?

- A. (1) and (2)
- B. (3) and (4)
- C. (3) and (6)
- D. (4) and (7)
- E. (6) and (7)

False dependencies

- We are still limited by **false dependencies**
- They are not “true” dependencies because they don’t have an arrow in data dependency graph
 - **WAR (Write After Read):** a later instruction overwrites the source of an earlier one
 - 2 and 1, 6 and 3, 7 and 4, 7 and 6
 - **WAW (Write After Write):** a later instruction overwrites the output of an earlier one
 - 6 and 1

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```



False dependencies

- We are still limited by **false dependencies**
- They are not “true” dependencies because they don’t have an arrow in data dependency graph


- **WAR (Write After Read):** a later instruction overwrites the source of an earlier one

- 2 and 1, 6 and 3, 7 and 4, 7 and 6

- **WAW (Write After Write):** a later instruction overwrites the output of an earlier one

- 6 and 1

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```



Announcements

- Assignment #3 due **this Thursday**
- Reading Quiz due next Tuesday

Computer Science & Engineering

203

つづく

