

Instruction Scheduling & Programming Modern Processors (I)

Hung-Wei Tseng

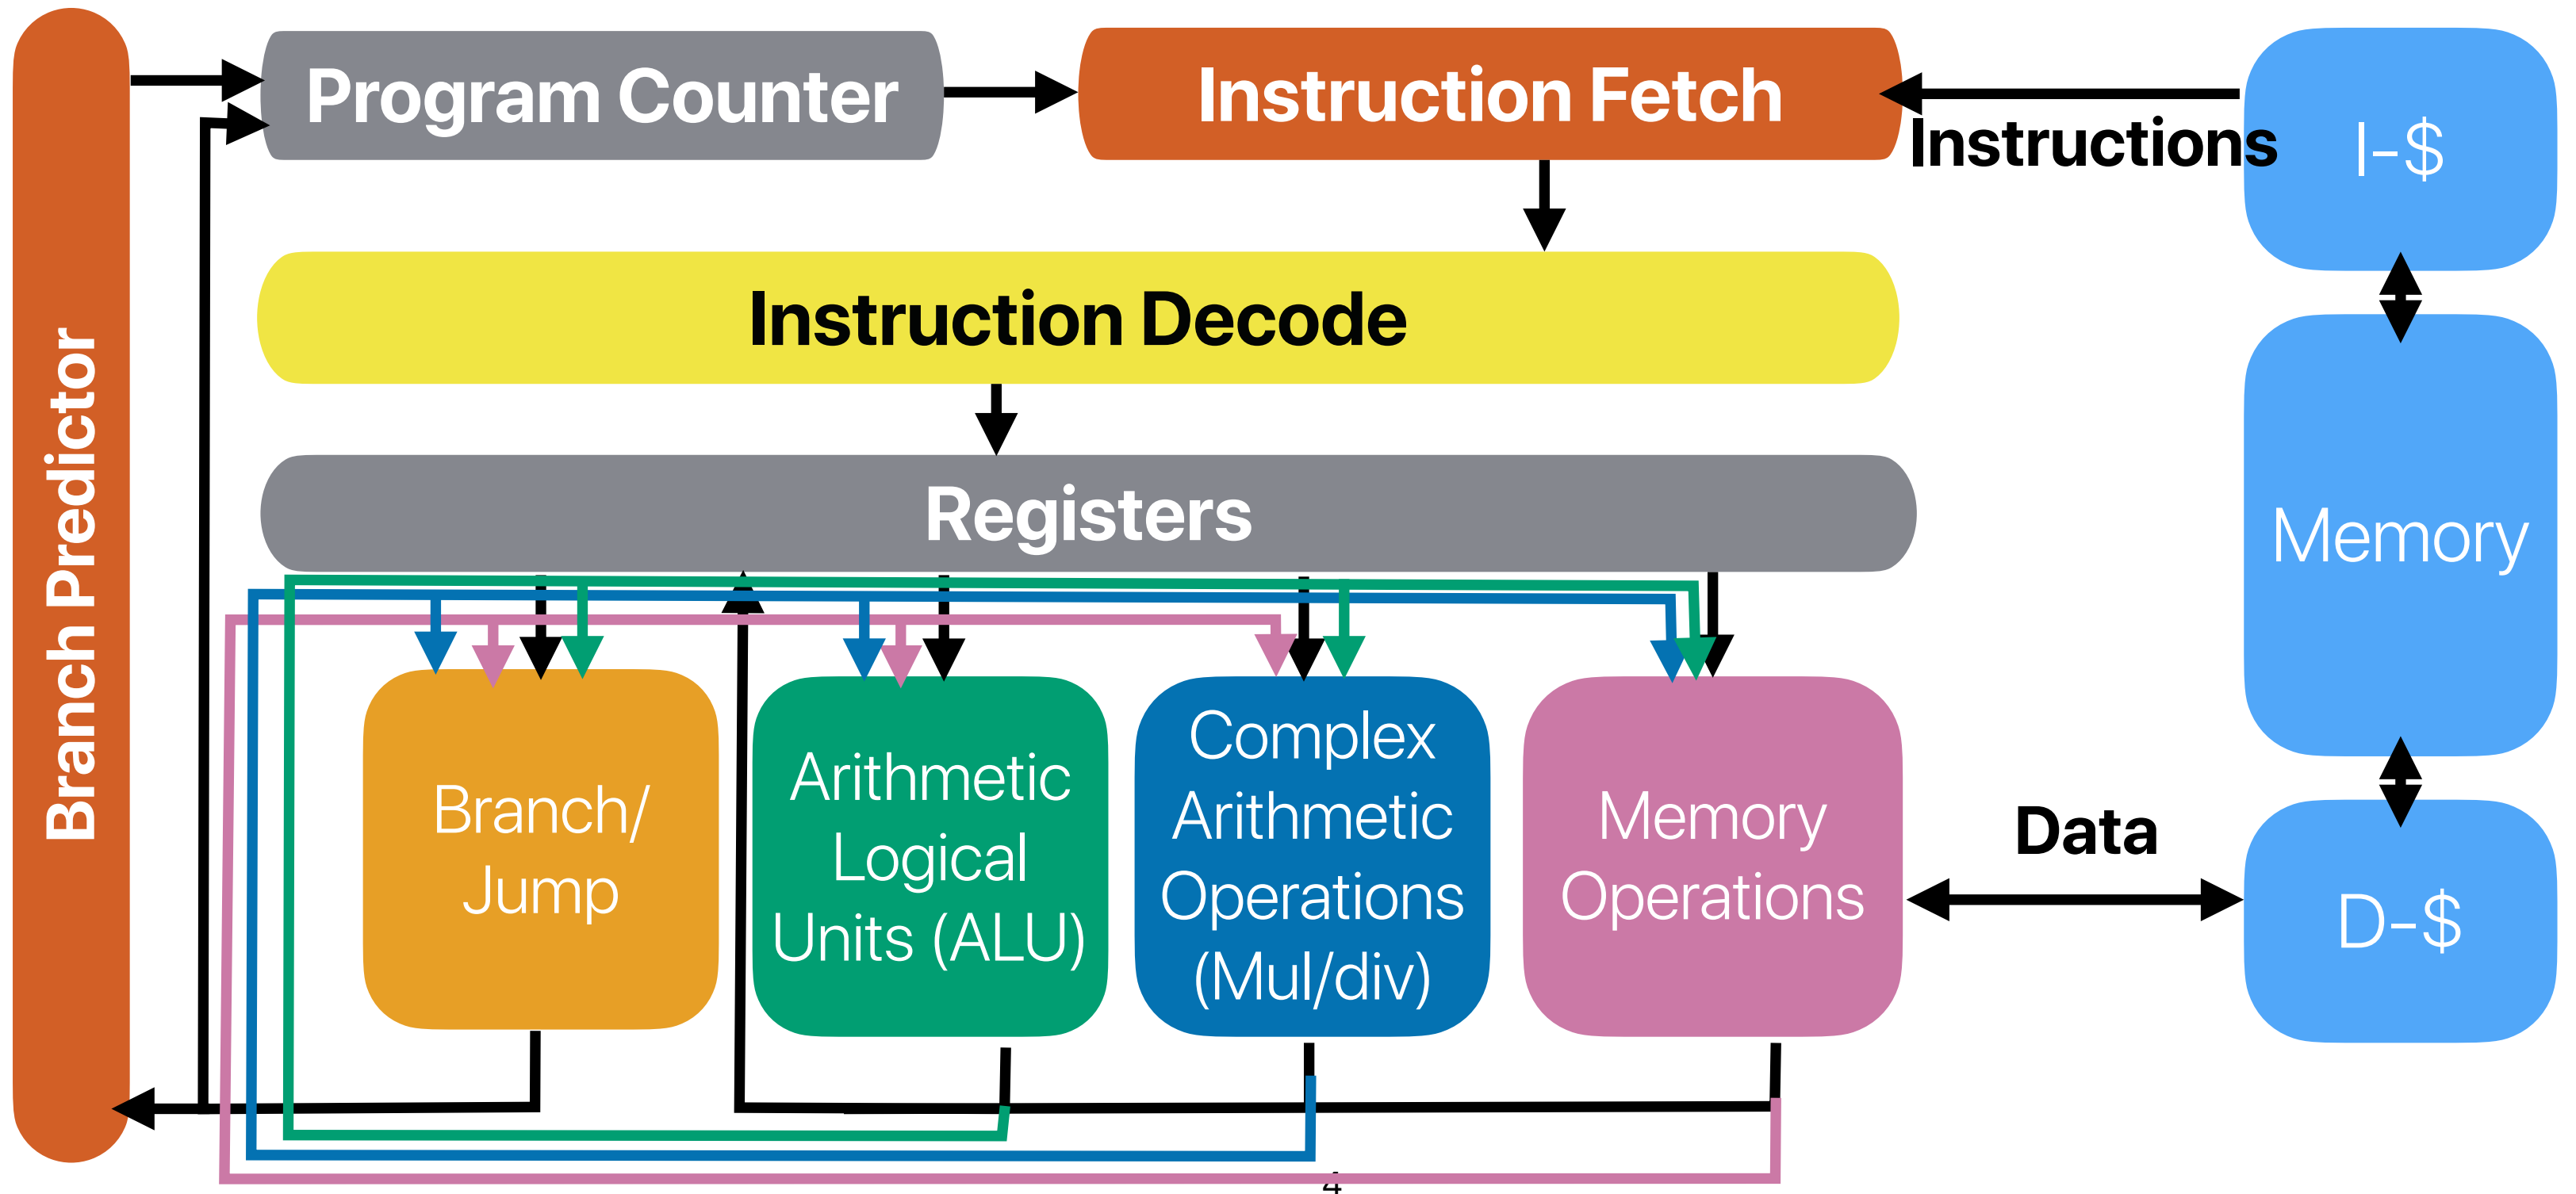
Recap: Three pipeline hazards

- Structural hazards — resource conflicts cannot support simultaneous execution of instructions in the pipeline
- Control hazards — the PC can be changed by an instruction in the pipeline
- Data hazards — an instruction depending on a the result that's not yet generated or propagated when the instruction needs that

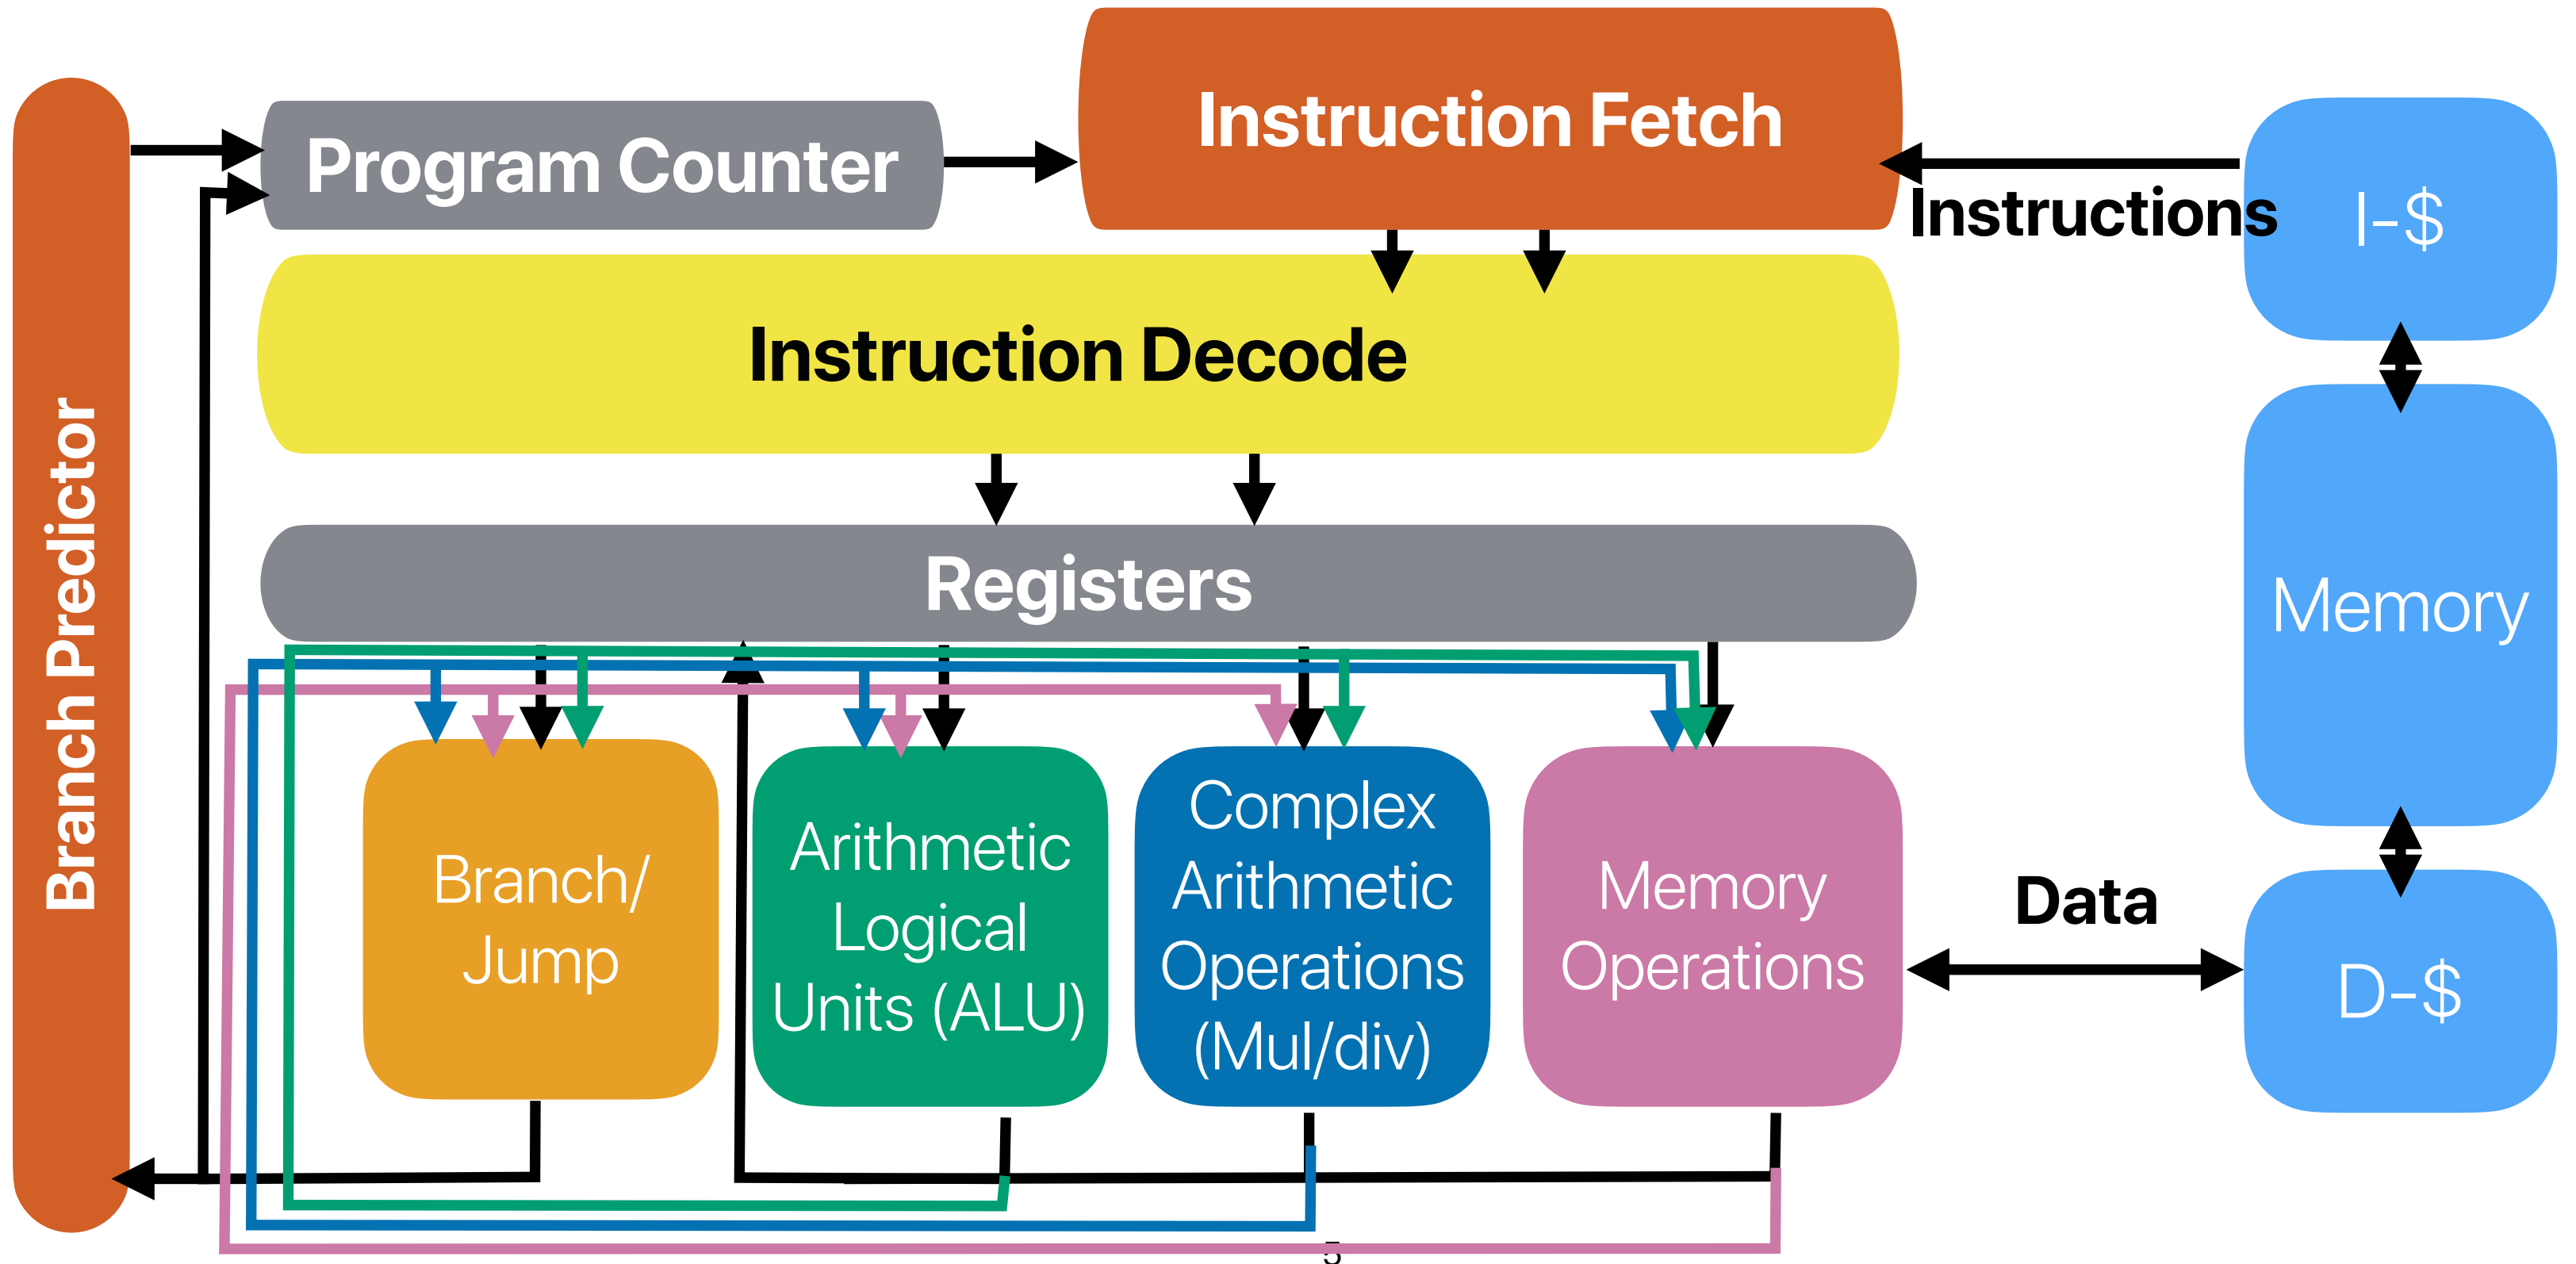
Recap: addressing hazards

- Structural hazards
 - Stall
 - Modify hardware design
- Control hazards
 - Stall
 - Static prediction
 - Dynamic prediction
- Data Hazards
 - Stall
 - Data forwarding
 - Dynamic instruction scheduling

Recap: Data "forwarding"



Super Scalar



Superscalar

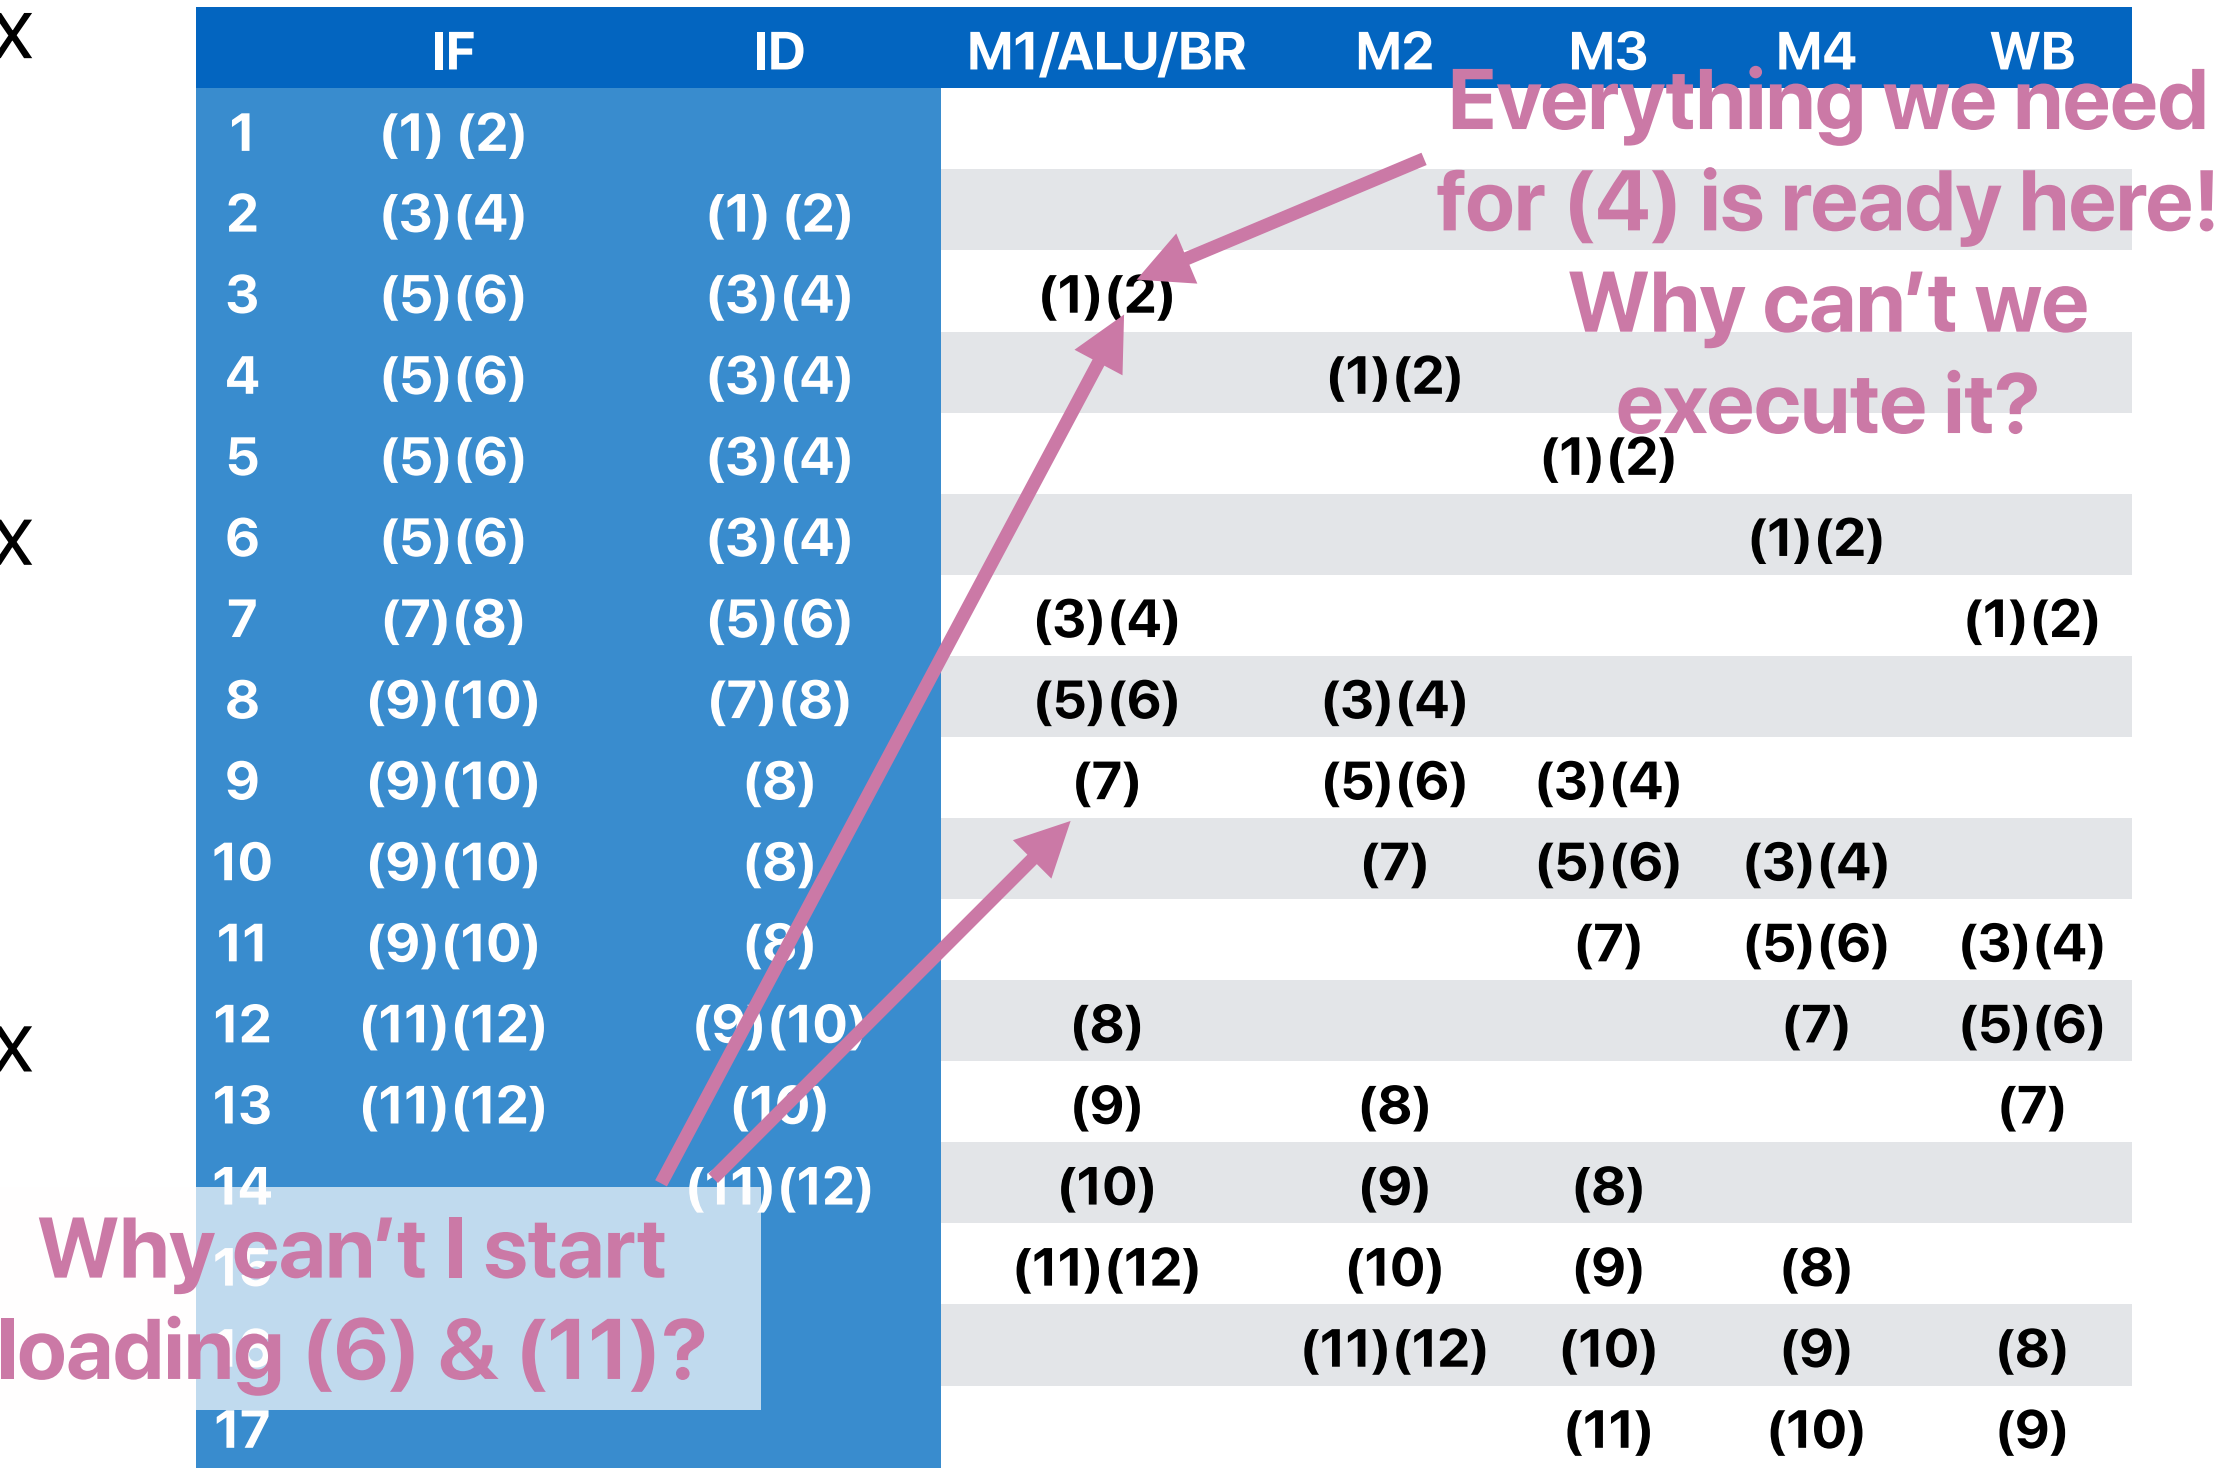
- Since we have many functional units now, we should fetch/decode more instructions each cycle so that we can have more instructions to issue!
- Super-scalar: fetch/decode/issue more than one instruction each cycle
 - **Fetch width:** how many instructions can the processor fetch/decode each cycle
 - **Issue width:** how many instructions can the processor issue each cycle
- The theoretical CPI should now be

1

$\min(\text{issue width}, \text{fetch width}, \text{decode width})$

If we loop many times (assume perfect predictor)

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
⑪ movl    (%rdi), %ecx
⑫ addq    $4, %rdi
⑬ addl    %ecx, %eax
⑭ cmpq    %rdx, %rdi
⑮ jne     .L3
```



False dependencies

- We are still limited by **false dependencies**
- They are not “true” dependencies because they don’t have an arrow in data dependency graph

- **WAR (Write After Read):** a later instruction overwrites the source of an earlier one


- 2 and 1, 6 and 3, 7 and 4, 7 and 6

- **WAW (Write After Write):** a later instruction overwrites the output of an earlier one

- 6 and 1

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```

8



Recap: Limitations of Compiler Optimizations

- If the hardware (e.g., pipeline changes), the same compiler optimization may not be that helpful
- The compiler can only optimize on static instructions, but cannot optimize dynamic instructions
- Compilers are limited by the registers an ISA provides

Outline

- Out-of-order, Dynamic instruction scheduling
- Programming on Modern Processor

Register renaming + speculative execution

- K. C. Yeager, "The Mips R10000 superscalar microprocessor," in IEEE Micro, vol. 16, no. 2, pp. 28-41, April 1996.

Recap: False dependencies

- We are still limited by **false dependencies**
- They are not “true” dependencies because they don’t have an arrow in data dependency graph

- **WAR (Write After Read):** a later instruction overwrites the source of an earlier one


- 2 and 1, 6 and 3, 7 and 4, 7 and 6

- **WAW (Write After Write):** a later instruction overwrites the output of an earlier one

We need to give each output a new register!!!


```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addq    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```

12



What if we can use more registers...

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```



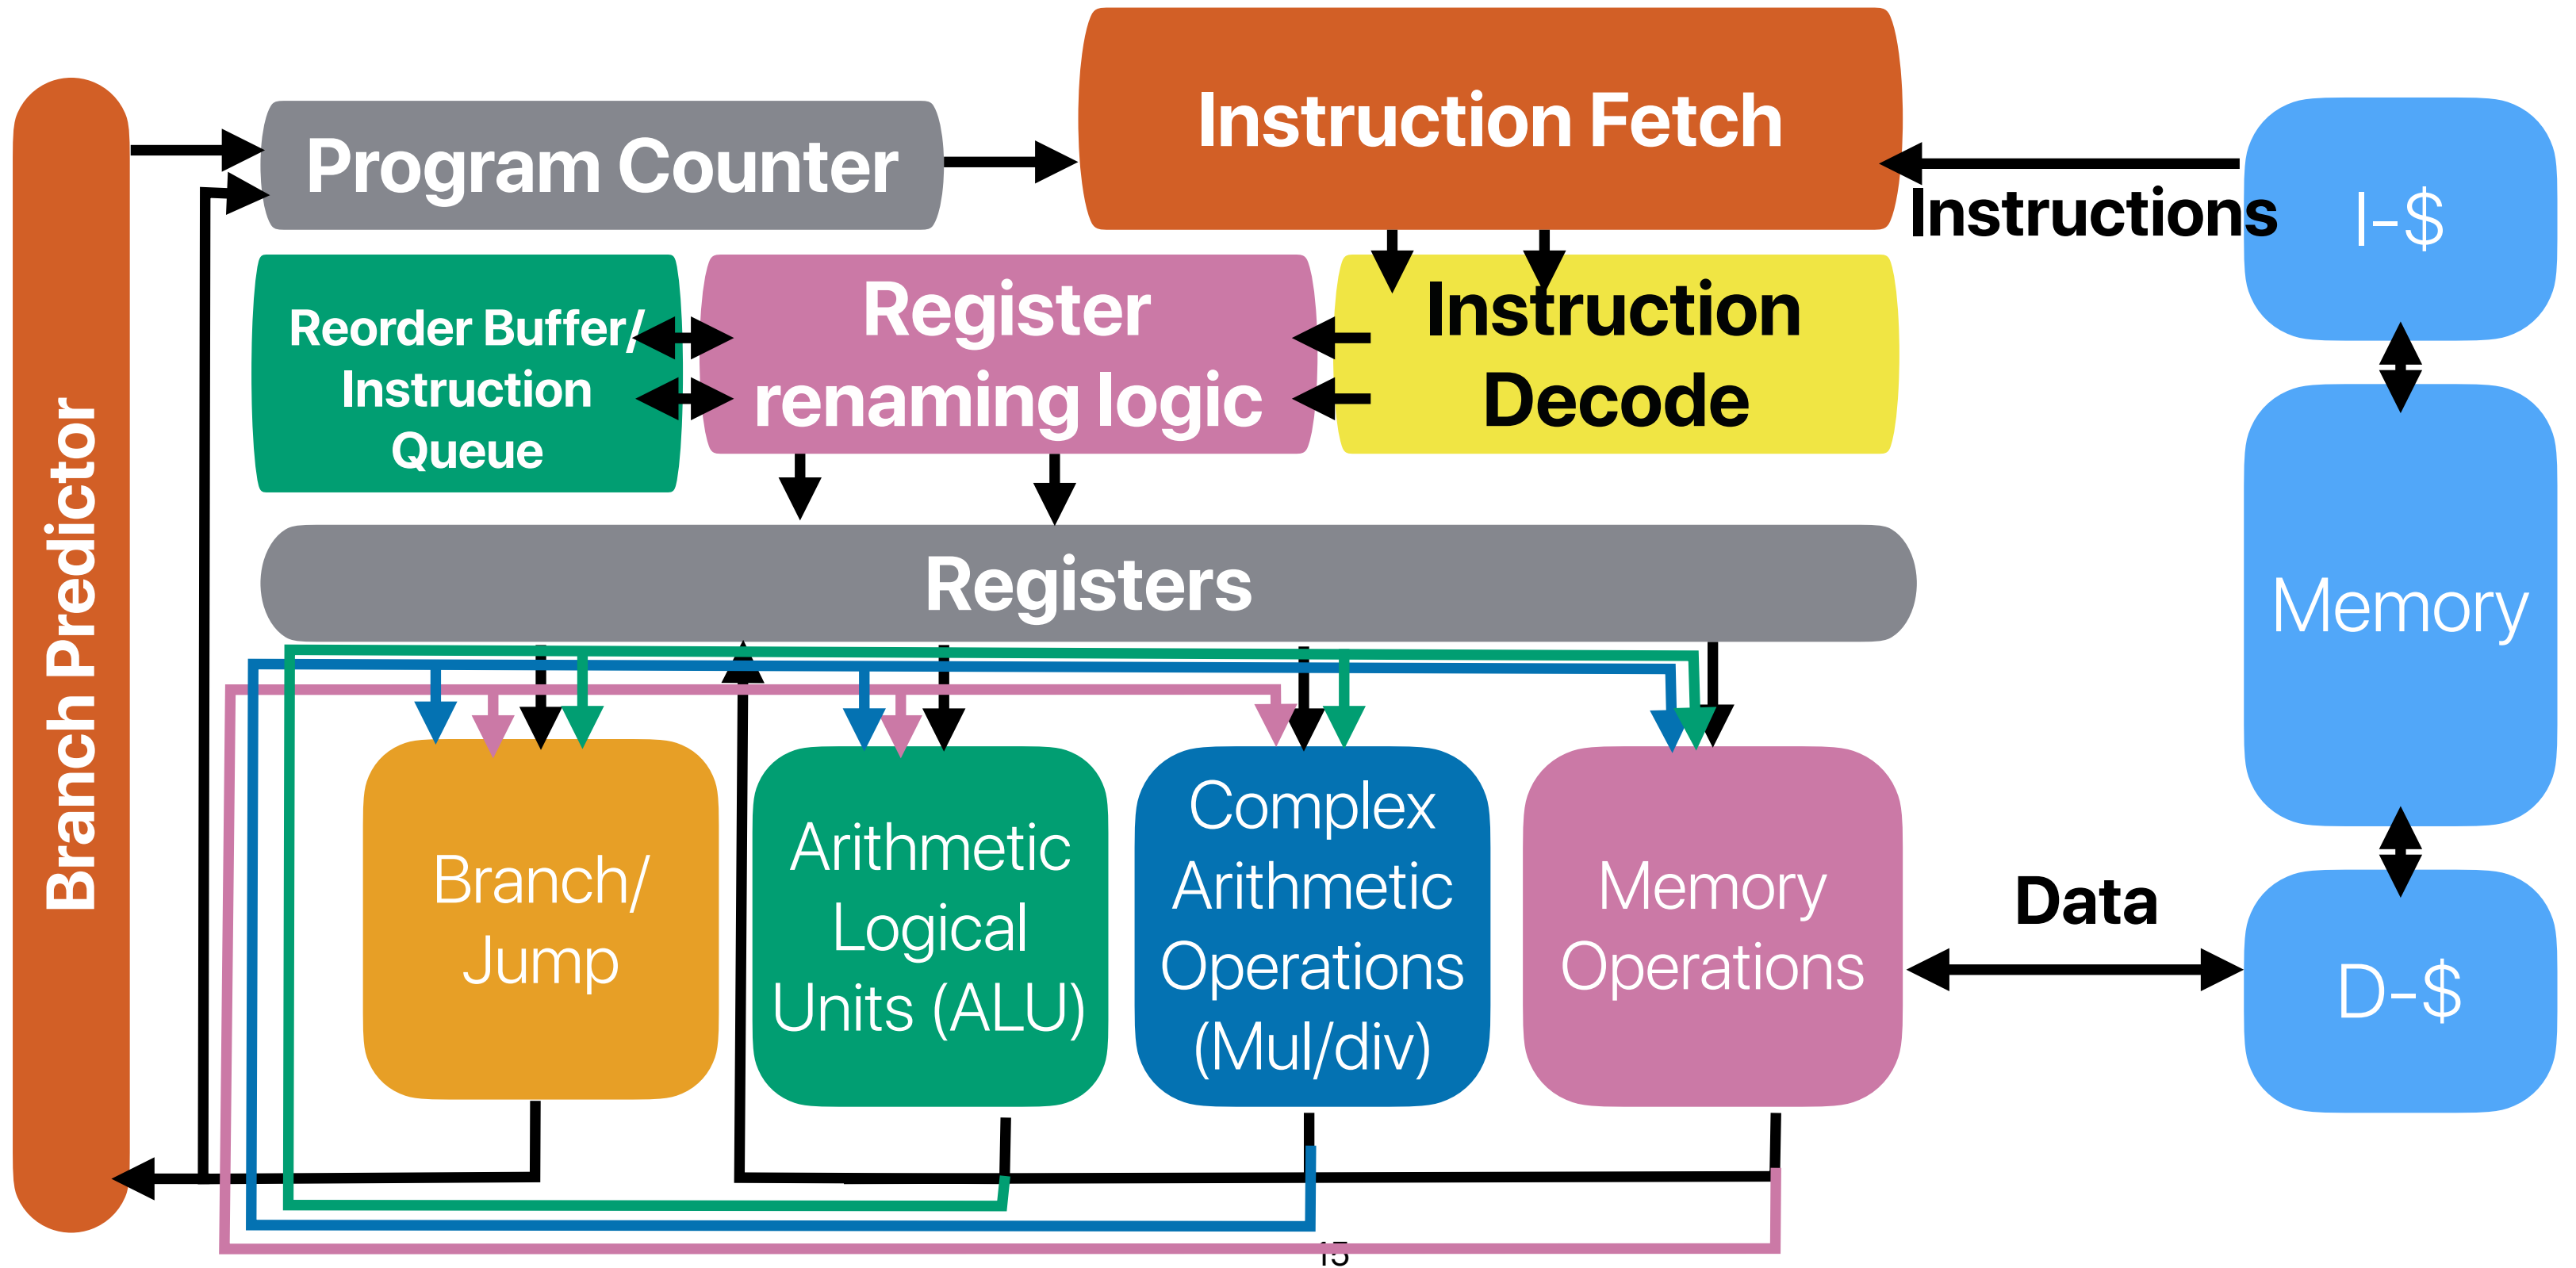
```
① movl    (%rdi), %ecx
② addq    $4, %rdi, %t0
③ addl    %ecx, %eax, %t1
④ cmpq    %rdx, %t0
⑤ jne     .L3
⑥ movl    (%t0), %t2
⑦ addq    $4, %t0, %t3
⑧ addl    %t1, %t2, %t4
⑨ cmpq    %rdx, %t3
⑩ jne     .L3
```

All false dependencies are gone!!!

Speculative Execution

- Exceptions (e.g. divided by 0, page fault) may occur anytime
 - A later instruction cannot write back its own result otherwise the architectural states won't be correct
- Hardware can schedule instruction across branch instructions with the help of branch prediction
 - Fetch instructions according to the branch prediction
 - However, branch predictor can never be perfect
- Execute instructions across branches
 - Speculative execution: execute an instruction before the processor know if we need to execute or not
 - Execute an instruction all operands are ready (the values of depending physical registers are generated)
 - Store results in **reorder buffer** before the processor knows if the instruction is going to be executed or not.

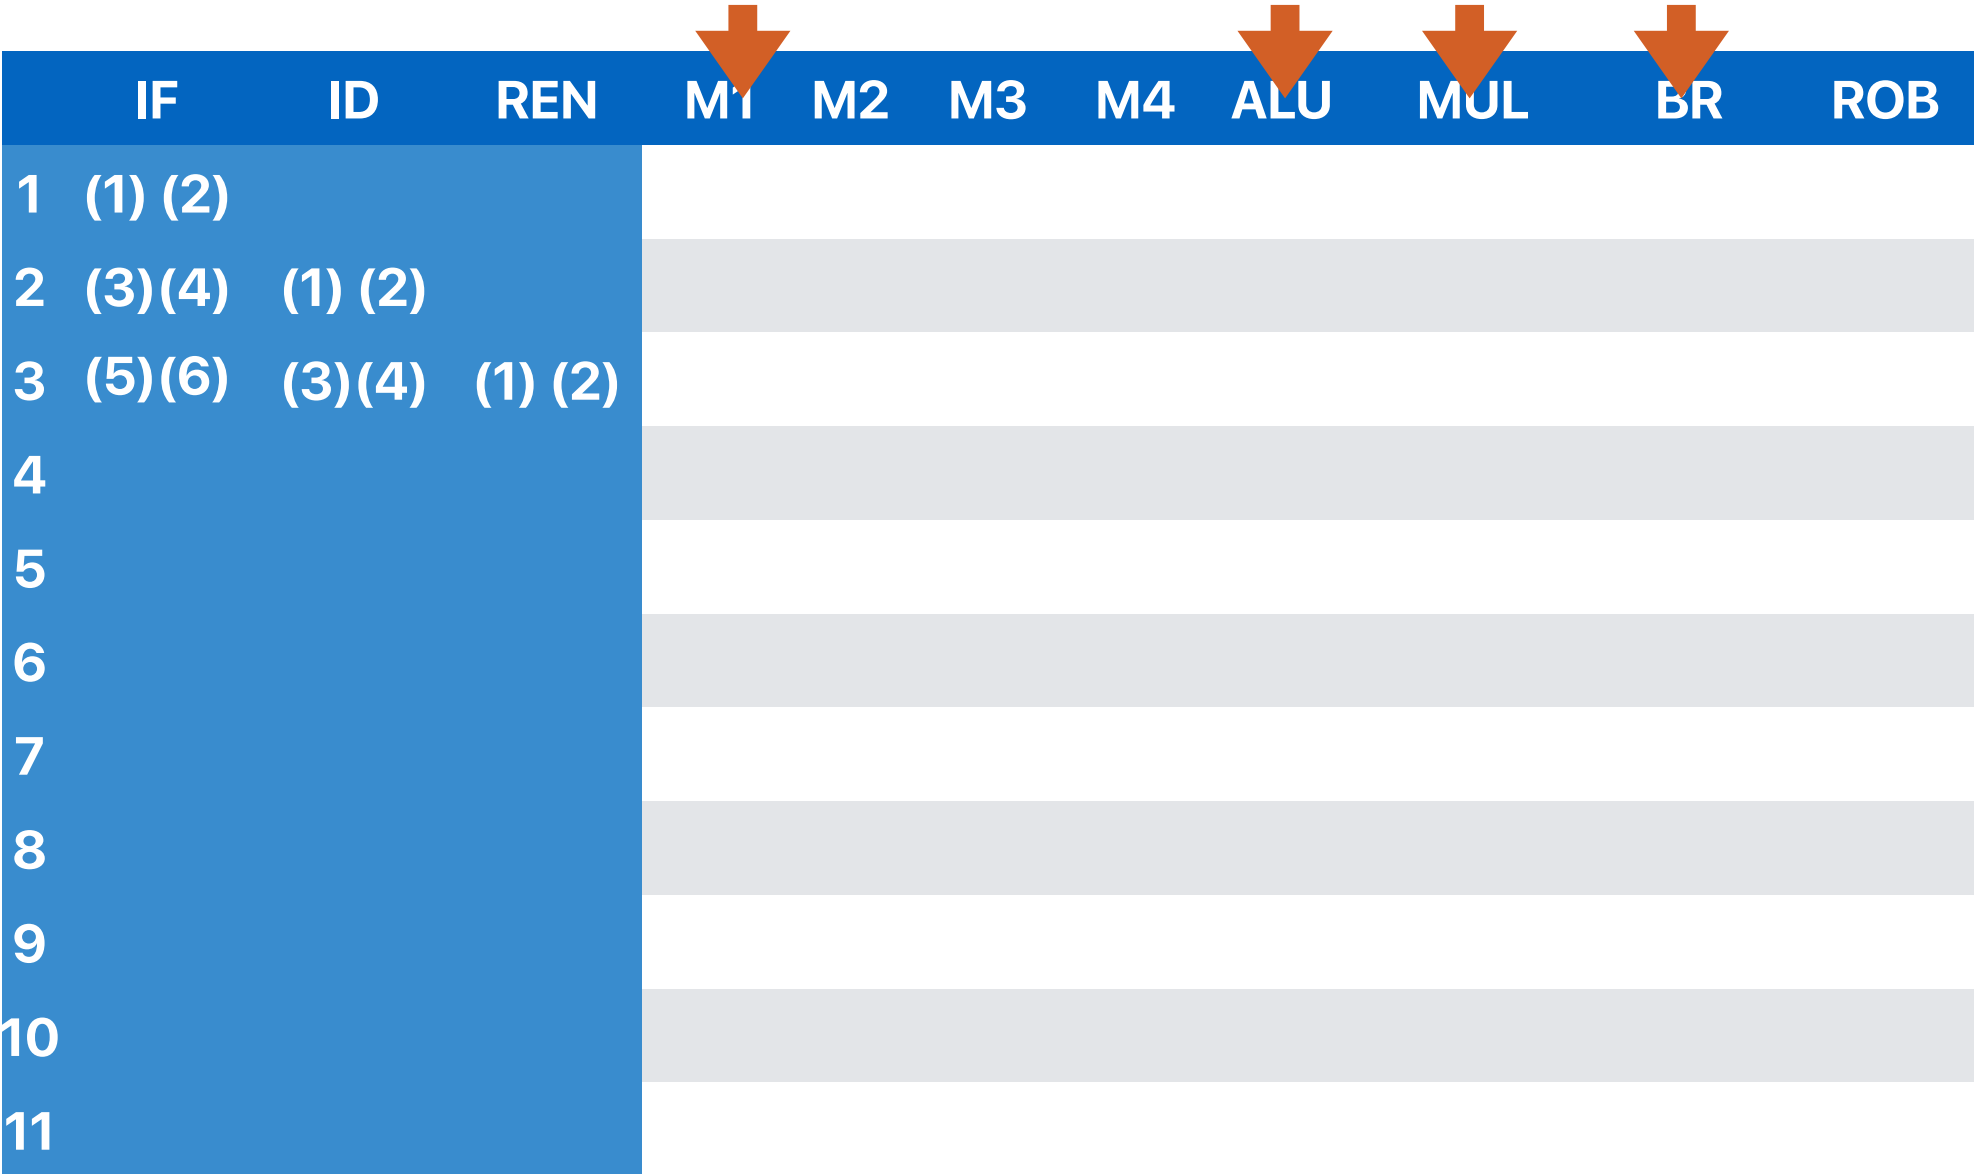
Register renaming



Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
⑪ movl    (%rdi), %ecx
⑫ addq    $4, %rdi
⑬ addl    %ecx, %eax
⑭ cmpq    %rdx, %rdi
⑮ jne     .L3
```



Physical Register	
eax	
ecx	
rdi	
rdx	

Valid	Value	In use	Valid	Value	In use
P1			P6		
P2			P7		
P3			P8		
P4			P9		
P5			P10		

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx
- ⑦

addq

\$4, %rdi
- ⑧

addl

%ecx, %eax
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx
- ⑫

addq

\$4, %rdi
- ⑬

addl

%ecx, %eax
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)	(4)	(1)	(2)							
3	(5)	(6)	(3)	(4)	(1)	(2)					
4		(5)	(6)	(3)	(4)						
5											
6											
7											
8											
9											
10											
11											

Physical Register	
eax	
ecx	P1
rdi	P2
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6			
P2	0		1	P7			
P3				P8			
P4				P9			
P5				P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

① movl (%rdi), %ecx → P1

② addq \$4, %rdi → P2

③ addl %ecx, %eax → P3

④ cmpq %rdx, %rdi

⑤ jne .L3

⑥ movl (%rdi), %ecx → P4

⑦ addq \$4, %rdi

⑧ addl %ecx, %eax

⑨ cmpq %rdx, %rdi

⑩ jne .L3


⑪ movl (%rdi), %ecx

⑫ addq \$4, %rdi

⑬ addl %ecx, %eax

⑭ cmpq %rdx, %rdi

⑮ jne .L3



(4) is now
executing before
(3)!

Physical Register	
eax	
ecx	P1
rdi	P2
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6			
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5				P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ① movl (%rdi), %ecx → P1

② addq \$4, %rdi → P2

③ addl %ecx, %eax → P3

④ cmpq %rdx, %rdi

⑤ jne .L3

⑥ movl (%rdi), %ecx → P4

⑦ addq \$4, %rdi → P5

⑧ addl %ecx, %eax → P6

⑨ cmpq %rdx, %rdi

⑩ jne .L3

⑪ movl (%rdi), %ecx

⑫ addq \$4, %rdi

⑬ addl %ecx, %eax

⑭ cmpq %rdx, %rdi

⑮ jne .L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)	(4)	(1)								
3	(5)	(6)	(3)								
4	(7)	(8)	(5)	(1)				(2)			
5	(9)	(10)	(7)		(1)			(4)			(2)
6	(11)	(12)	(9)	(6)		(1)				(5)	(2)(4)
7											
8											
9											
10											
11											

Physical Register	
eax	P6
ecx	P1
rdi	P5
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5	0		1	P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx
- ⑫

addq

\$4, %rdi
- ⑬

addl

%ecx, %eax
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8											
9											
10											
11											

Physical Register	
eax	P6
ecx	P1
rdi	P5
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5	0		1	P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9											
10											
11											

Physical Register	
eax	P6
ecx	P7
rdi	P8
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	1		1	P6	0		1
P2	1		1	P7	0		1
P3	0		1	P8	0		1
P4	0		1	P9			
P5	1		1	P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)	(4)	(1)								
3	(5)	(6)	(3)								
4	(7)	(8)	(5)	(1)				(2)			
5	(9)	(10)	(7)		(1)			(4)			(2)
6	(11)	(12)	(9)	(6)		(1)				(5)	(2)(4)
7	(13)	(14)	(11)		(6)		(1)	(7)			(2)(4)(5)
8	(15)	(16)	(13)			(6)		(3)			(1)(2)(4)(5) (7)
9	(17)	(18)	(15)	(11)			(6)	(9)			(3)(4)(5)(7)
10											
11											

Physical Register	
eax	P9
ecx	P7
rdi	P8
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	1		0	P6	0		1
P2	1		0	P7	0		1
P3	1		1	P8	0		1
P4	0		1	P9	0		1
P5	1		1	P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11											

Physical Register	
eax	P9
ecx	P7
rdi	P8
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	1		0	P6	0		1
P2	1		0	P7	0		1
P3	1		0	P8	0		1
P4	0		1	P9	0		1
P5	1		1	P10	0		1

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11		(19)(20)	(13)(14)(15)(16)(17)(18)			(11)		(12)			(8)(9)(10)

Physical Register	
eax	P6
ecx	P1
rdi	P5
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	1		0	P6	1		1
P2	1		0	P7	0		1
P3	1		0	P8	0		1
P4	1		0	P9	0		1
P5	1		1	P10	0		1

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11		(19)(20)	(13)(14)(15)(16)(17)(18)			(11)		(12)			(8)(9)(10)
12			(13)(15)(17)(18)(19)(20)	(16)			(11)	(14)			(12)
13											
14											
15											

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

① movl (%rdi), %ecx → P1

② addq \$4, %rdi → P2

③ addl %ecx, %eax → P3

④ cmpq %rdx, %rdi

⑤ jne .L3

⑥ movl (%rdi), %ecx → P4

⑦ addq \$4, %rdi → P5

⑧ addl %ecx, %eax → P6

⑨ cmpq %rdx, %rdi

⑩ jne .L3

⑪ movl (%rdi), %ecx → P7

⑫ addq \$4, %rdi → P8

⑬ addl %ecx, %eax → P9

⑭ cmpq %rdx, %rdi

⑮ jne .L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)	(4)	(1)								
3	(5)	(6)	(3)								
4	(7)	(8)	(5)	(1)				(2)			
5	(9)	(10)	(7)		(1)			(4)			(2)
6	(11)	(12)	(9)	(6)		(1)				(5)	(2)(4)
7	(13)	(14)	(11)		(6)		(1)	(7)			(2)(4)(5)
8	(15)	(16)	(13)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)	(18)	(15)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)	(20)	(17)		(11)			(8)		(10)	(6)(7)(9)
11		(19)	(20)			(11)		(12)			(8)(9)(10)
12			(13)	(16)			(11)	(14)			(12)
13			(17)		(16)			(13)		(15)	(11)(12)(14)
14											
15											

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11		(19)(20)	(13)(14)(15)(16)(17)(18)			(11)		(12)			(8)(9)(10)
12			(13)(15)(17)(18)(19)(20)	(16)			(11)	(14)			(12)
13			(17)(18)(19)(20)		(16)			(13)		(15)	(11)(12) (14)
14						(16)		(17)			(13)(14)(15)
15											

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

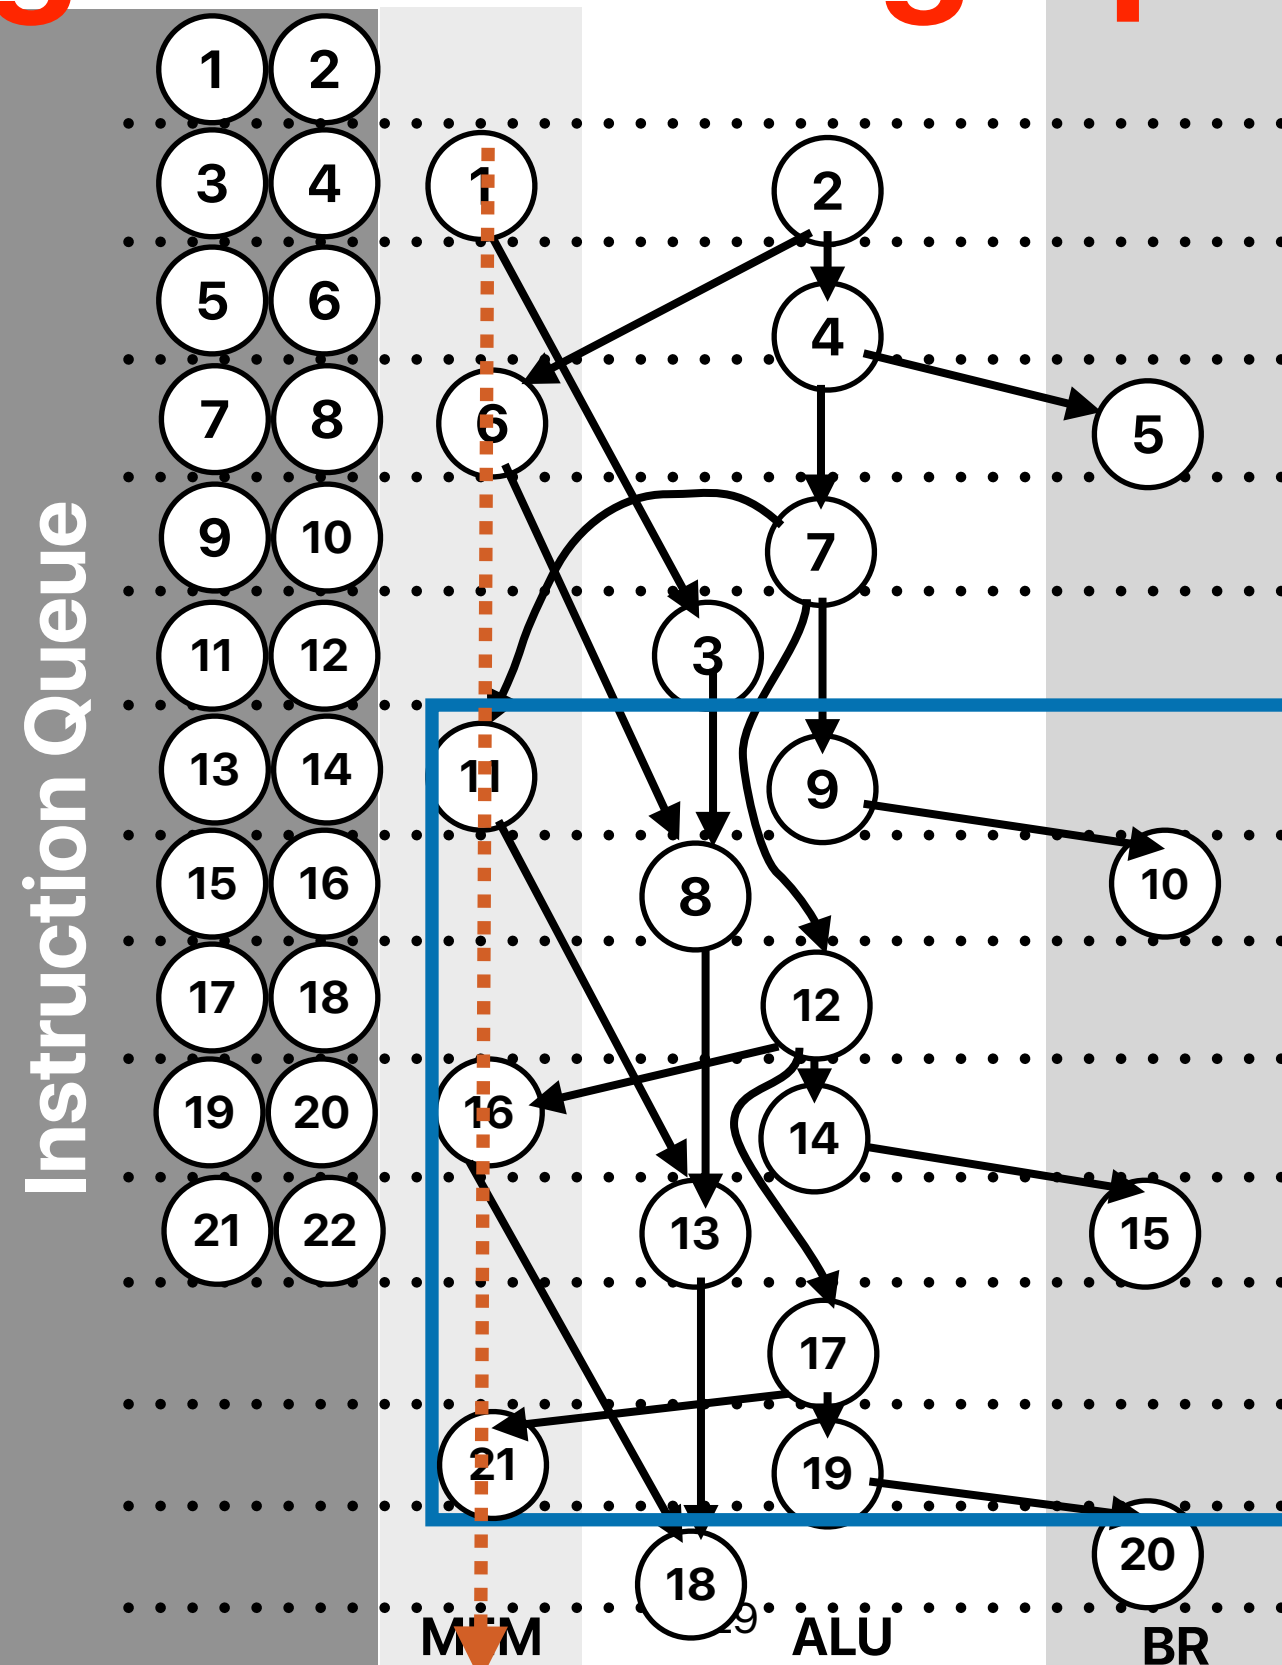
jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11		(19)(20)	(13)(14)(15)(16)(17)(18)			(11)		(12)			(8)(9)(10)
12			(13)(15)(17)(18)(19)(20)	(16)			(11)	(14)			(12)
13			(17)(18)(19)(20)		(16)			(13)		(15)	(11)(12)(14)
14						(16)		(17)			(13)(14)(15)
15							(16)	(19)			(17)

Through data flow graph analysis

```
① movl (%rdi), %ecx
② addq $4, %rdi
③ addl %ecx, %eax
④ cmpq %rdx, %rdi
⑤ jne .L3
⑥ movl (%rdi), %ecx
⑦ addq $4, %rdi
⑧ addl %ecx, %eax
⑨ cmpq %rdx, %rdi
⑩ jne .L3
⑪ movl (%rdi), %ecx
⑫ addq $4, %rdi
⑬ addl %ecx, %eax
⑭ cmpq %rdx, %rdi
⑮ jne .L3
⑯ movl (%rdi), %ecx
⑰ addq $4, %rdi
⑱ addl %ecx, %eax
⑲ cmpq %rdx, %rdi
⑳ jne .L3
㉑ movl (%rdi), %ecx
```



Execution time is determined
by the "critical path"
composed by 1, 6, 11, ..., 1+5n


3 cycles every iteration
 $CPI = \frac{3}{5} = 0.6!$

What about "linked list"

- Assume the current PC is already at instruction (1) and this linked list has only three nodes. This processor can fetch and issue 2 instructions per cycle, with exactly the same register renaming hardware and pipeline as we showed previously.

Which of the following C state of the code snippet determines the performance?

- A. `do {`
- B. `number_of_nodes++;`
- C. `current = current->next;`
- D. `} while (current != NULL);`



What about "linked list"

Performance determined by the critical path


4 cycles each iteration

4 instructions per iteration

$$CPI = \frac{4}{4} = 1$$


```
do {  
    number_of_nodes++;  
    current = current->next;  
} while ( current != NULL );
```

①	.L3:	movq	8(%rdi), %rdi
②		addl	\$1, %eax
③		testq	%rdi, %rdi
④		jne	.L3



The pipelines of Modern Processors

Intel Skylake



Recap: Intel Skylake




Figure 4. Skylake core block diagram.


Intel Alder Lake

$$MinCPI = \frac{1}{12}$$

$$MinINTInst . CPI = \frac{1}{5}$$

$$MinMEMInst . CPI = \frac{1}{7}$$

$$MinBRInst . CPI = \frac{1}{2}$$



AMD Zen 3 (RyZen 5000 Series)

3-issue memory pipeline

4-issue integer pipeline + 1 additional branch

$$MinCPI = \frac{1}{8}$$

$$MinINTInst.CPI = \frac{1}{4}$$

$$MinMEMInst.CPI = \frac{1}{3}$$

$$MinBRInst.CPI = \frac{1}{2}$$




FIGURE 1. "Zen 3" block diagram.

Summary: Characteristics of modern processor architectures

- Multiple-issue pipelines with multiple functional units available
 - Multiple ALUs
 - Multiple Load/store units
 - Dynamic OoO scheduling to reorder instructions whenever possible
- Cache
- Branch predictors

Performance Programming on Modern Processors

Demo: Popcount

- The population count (or popcount) of a specific value is the number of set bits (i.e., bits in 1s) in that value.
- Applications
 - Parity bits in error correction/detection code
 - Cryptography
 - Sparse matrix
 - Molecular Fingerprinting
 - Implementation of some succinct data structures like bit vectors and wavelet trees.

Demo: pop count

- Given a 64-bit integer number, find the number of 1s in its binary representation.

- Example 1:

Input: 9487

Output: 7

Explanation: 9487's binary representation is
0b10010100001111

```
int main(int argc, char *argv[]) {  
  
    uint64_t key = 0xdeadbeef;  
  
    int count = 1000000000;  
    uint64_t sum = 0;  
  
    for (int i=0; i < count; i++)  
    {  
        sum += popcount(RandLFSR(key));  
    }  
    printf("Result: %lu\n", sum);  
    return sum;  
}
```


Five implementations

- Which of the following implementations will perform the best on modern pipeline processors?

A

```
inline int popcount(uint64_t x){
    int c=0;
    while(x) {
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```

B

```
inline int popcount(uint64_t x) {
    int c = 0;
    while(x) {
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```

C

```
inline int popcount(uint64_t x) {
    int c = 0;
    int table[16] = {0, 1, 1, 2, 1,
2, 2, 3, 1, 2, 2, 3, 2, 3, 3, 4};
    while(x) {
        c += table[(x & 0xF)];
        x = x >> 4;
    }
    return c;
}
```

D

```
inline int popcount(uint64_t x) {
    int c = 0;
    int table[16] = {0, 1, 1, 2, 1,
2, 2, 3, 1, 2, 2, 3, 2, 3, 3, 4};
    for (uint64_t i = 0; i < 16; i++)
    {
        c += table[(x & 0xF)];
        x = x >> 4;
    }
    return c;
}
```

E

```
inline int popcount(uint64_t x) {
    int c = 0;
    for (uint64_t i = 0; i < 16; i++)
    {
        switch((x & 0xF))
        {
            case 1: c+=1; break;
            case 2: c+=1; break;
            case 3: c+=2; break;
            case 4: c+=1; break;
            case 5: c+=2; break;
            case 6: c+=2; break;
            case 7: c+=3; break;
            case 8: c+=1; break;
            case 9: c+=2; break;
            case 10: c+=2; break;
            case 11: c+=3; break;
            case 12: c+=2; break;
            case 13: c+=3; break;
            case 14: c+=3; break;
            case 15: c+=4; break;
            default: break;
        }
        x = x >> 4;
    }
    return c;
}
```

B

C

D

Five implementations

- Which of the following implementations will perform the best on modern pipeline processors?

A

```
inline int popcount(uint64_t x){
    int c=0;
    while(x) {
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```

B

```
inline int popcount(uint64_t x) {
    int c = 0;
    while(x) {
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```

C

```
inline int popcount(uint64_t x) {
    int c = 0;
    int table[16] = {0, 1, 1, 2, 1, 2, 2, 3, 1, 2, 2, 3, 2, 3, 3, 4};
    while(x) {
        c += table[(x & 0xF)];
        x = x >> 4;
    }
    return c;
}
```

D

```
inline int popcount(uint64_t x) {
    int c = 0;
    int table[16] = {0, 1, 1, 2, 1, 2, 2, 3, 1, 2, 2, 3, 2, 3, 3, 4};
    for (uint64_t i = 0; i < 16; i++) {
        c += table[(x & 0xF)];
        x = x >> 4;
    }
    return c;
}
```

E

```
inline int popcount(uint64_t x) {
    int c = 0;
    for (uint64_t i = 0; i < 16; i++) {
        switch((x & 0xF)) {
            case 1: c+=1; break;
            case 2: c+=1; break;
            case 3: c+=2; break;
            case 4: c+=1; break;
            case 5: c+=2; break;
            case 6: c+=2; break;
            case 7: c+=3; break;
            case 8: c+=1; break;
            case 9: c+=2; break;
            case 10: c+=2; break;
            case 11: c+=3; break;
            case 12: c+=2; break;
            case 13: c+=3; break;
            case 14: c+=3; break;
            case 15: c+=4; break;
            default: break;
        }
        x = x >> 4;
    }
    return c;
}
```

Announcements

- Assignment #3 due **Friday**
- Reading Quiz due next Monday
- iEval
 - Submit by 6/9 and take a screenshot of your submission
 - Submit your screen screenshot in gradescope — it counts as a “full-credit” notebook assignment (technically help to drop two of your lowest notebook assignments)

Computer Science & Engineering

203

つづく

